

Exercise 1:

Course name : .Build a Convolution Neural Network for Image Recognition.

Go through the modules of the course mentioned and answer the self-assessment questions given in the link below at the end of the course.

Representation of an image

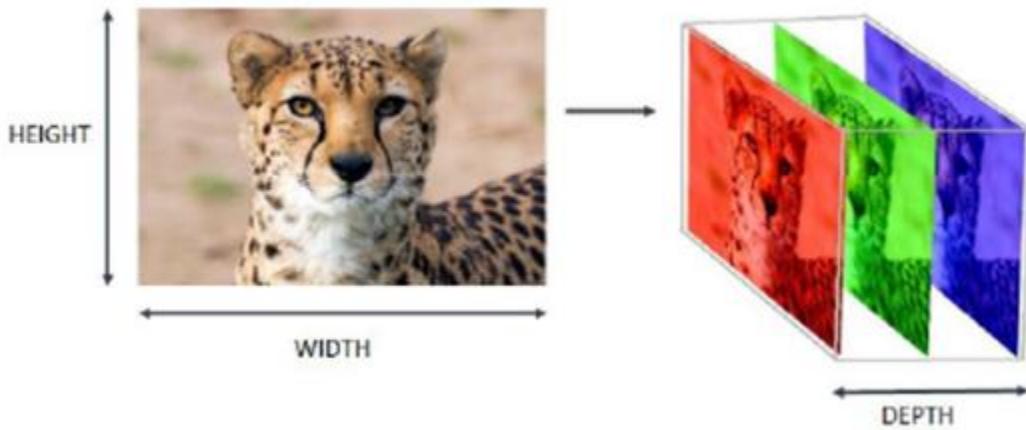
An image is an unstructured data represented by pixels.

To store an image on a computer, the image is broken down into tiny elements called pixels (short for picture element). A pixel is a number and it represents one colour. Usually, an 8 bit number is used to represent a pixel. Hence, each pixel of a black and white image can have $2^8 = 256$ different shades of gray. It can also be represented as a decimal number from 0 to 255. Zero stands for black and 255 for white.

An image with a resolution of 1024 by 798 pixels has 1024 x 798 total number of pixels (817,152 pixels)

A pixel in a coloured image is an array of 3 numbers. They represent the red, green and blue (RGB) values. For an 8 bit representation, each colour can take 256 shades. The final colour of the pixel is given by the combination of the 3 values. Hence each pixel can have $256 \times 256 \times 256$ values of nearly 17 million different colours.

So, the computer sees an image as an array of numbers. It represents an image with a colour depth that represents the number of bits that represent a pixel (in the examples mentioned above, we considered it to be 8) and the image resolution: the height and width of an image in pixels. As a colour image has 3 matrices, one for each colour, hence it is represented by its height, width and depth. It is sometimes called a 3D volume.



This figure shows the RGB decomposition of a colour image.

Feature

Can we teach a machine to identify objects the way humans do? Assume that we have a large number of images of cats. Is there a way to train a machine to identify cats with these given images? Can the machine identify a new cat that was not in the training dataset?

To answer these questions, we have to understand how we identify objects. We identify objects by their



features. Lets understand what is a feature.

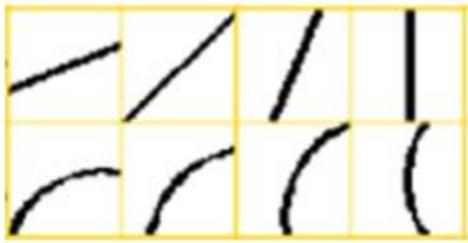
Consider the two images a caracal and a cheetah:

- Caracal has pointy ears. Cheetah has round ears.
- Cheetah is spotted. Caracal is not spotted.
- Cheetah has dark black lines on its face.

Such differentiating data is called feature.

The abstraction of such features (a straight line, a curve , etc.) is shown in the diagram. A feature is a small image, usually of size 2×2 or 3×3 pixel dimension. These features are also called filters, kernels or weights in Deep learning.

0.5	0.2	0.1
0.2	0.5	0.2
0.1	0.2	0.5



But, how does a machine find out what features an image has? It learns features using the mathematical operation called convolution. Convolution means the similarity of a feature with a given image. We select a random feature and convolve (roll) it on the surface of the image. The learning process progressively modifies the random feature such that it has a good match with the given image. Multiple features are deployed, each starting at a different random state. Each of them specializes to learn different features of the image. The process of learning features from an image is called feature extraction.

Once we have extracted the features from an object, we can use them to identify the object.

Learnable filters

Convolutional neural network has two parts. The first part identifies the features in the image. The features learnt in the first phase is fed to the second part, a fully connected neural network. This part does the classification.

The question is how do we select a filter such that its overlap with the image is maximum? Convolutional neural networks uses learnable filters. We randomly select such filters and then learn the required filters by training, just as in ANN, where we begin with random weights and then learn the best weights through training. This is why filters are also called weights.

Sample filter -

0.5	0.2	0.1
0.2	0.5	0.2
0.1	0.2	0.5

Padding

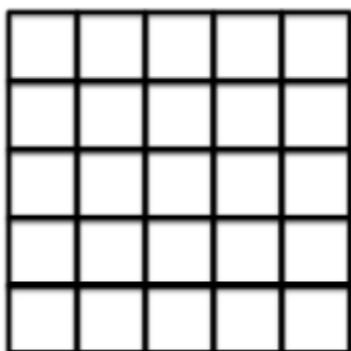
We observed that everytime we use the filter and scan the image for similarity, the size of the resulting image becomes smaller and smaller. We do not want that, because we want to preserve the original size of the image to extract some low level of features. The second issue is that the pixels in the periphery get covered only once but the ones in the middle get covered multiple times.

Thus, there are two issues with convolution -

- Shrinking feature map and
- Losing information on the periphery of the image.

These two issues are solved by padding. In padding, the image is added with extra pixels at the boundary before the convolution. Zero padding means the value of these added pixels is zero. Sometimes the value of the edge pixel is added in the padded cell.

Add image



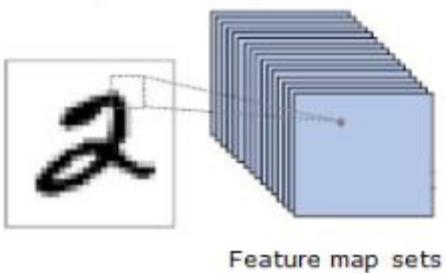
If zero padding = 1 , one pixel will be added around the original image with value = 0.

0	0	0	0	0	0	0
0						0
0						0
0						0
0						0
0						0
0	0	0	0	0	0	0

Feature map sets

As discussed, the feature map is a measure of similarity of the filter with the image. An image can have multiple distinguishing features. To correctly identify the image we need to learn all of them. Hence we deploy a number of learnable filters. Each specializes to recognize one feature. And each filter results in a feature map. Hence we get a set for feature maps.

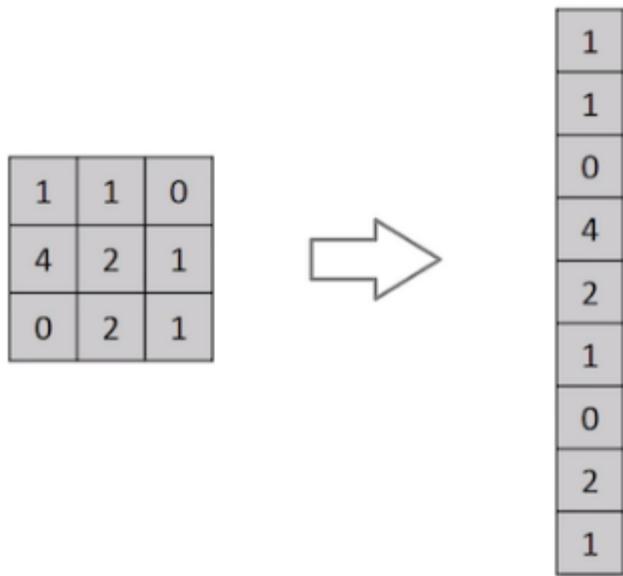
Part drawing showing feature map sets



Flattening

How are the pooled layer sets fed to a fully connected layer? Each element from the pooled set is still a matrix. However, a fully connected network does not accept a matrix but only scalars.

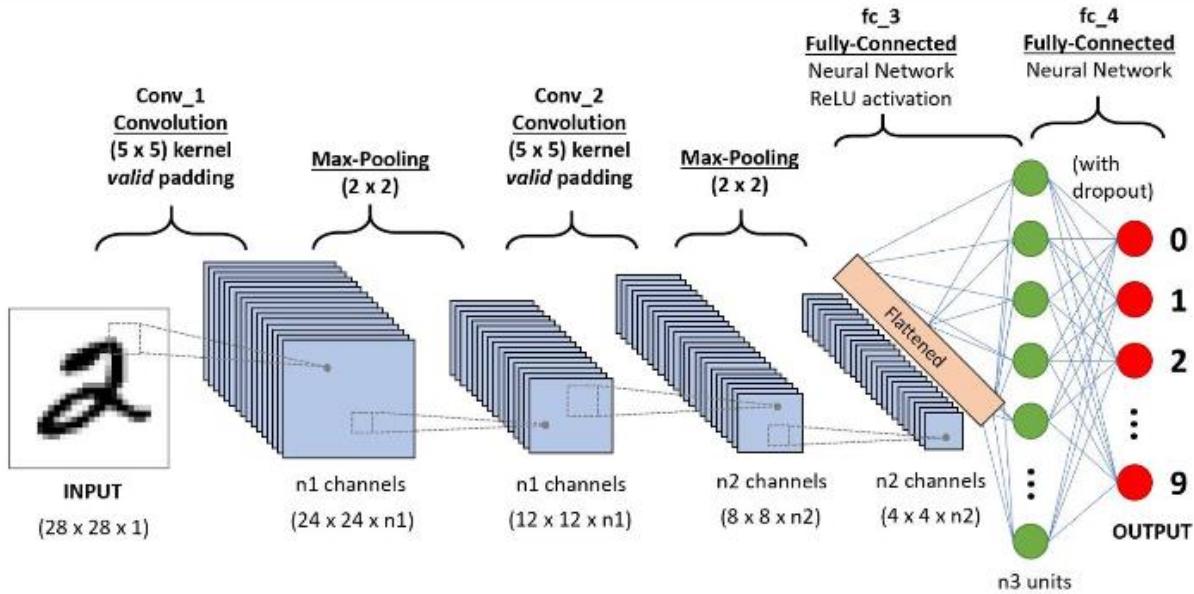
So we flatten the matrix; See the image below.\



Flattening of a 3x3 image matrix into a 9x1 vector

Neural networks can very easily overfit -- they learn very well on training set but perform poorly on validation set.

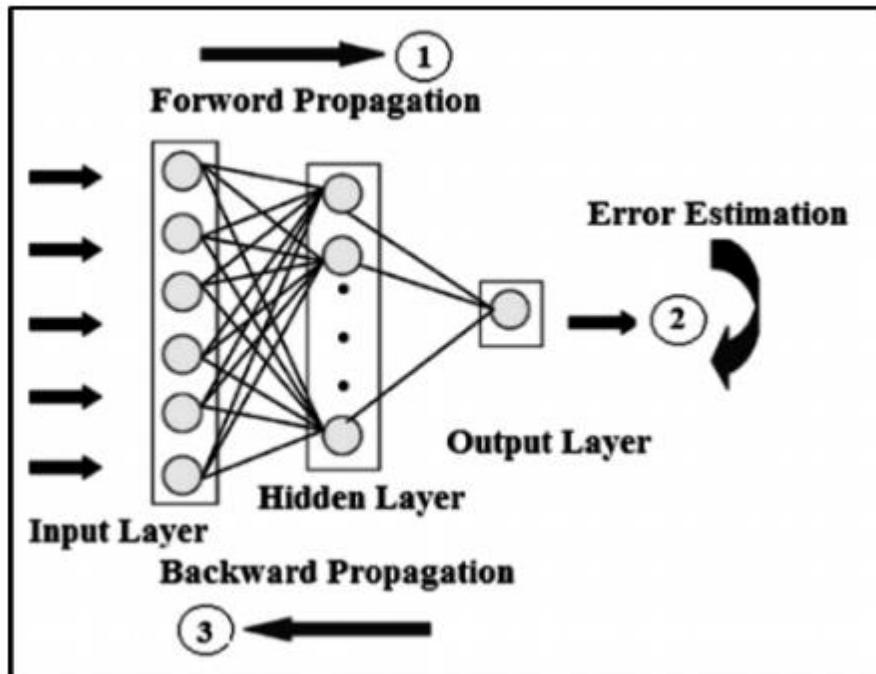
To get around this problem, the network is trained by randomly dropping a few weights (or filters) during each pass. This ensures that the filters do not “over specialize” in recognizing features and can work when a few features are marked off.



A CNN sequence to classify handwritten digits

Training

The fully connected network at the end is trained using Back-propagation and gradient descent discussed earlier. It is during this process that the distinguishing features of the images are learnt by the filters.



Exercise 2

Module name: Understanding and Using ANN : Identifying age group of an actor Exercise : Design Artificial Neural Networks for Identifying and Classifying an actor using Kaggle Dataset.

We are going to take a scenario of identifying the age group of various movie characters just by considering their facial attributes and in turn will try to understand the implementation of deep neural networks in python.

We will use the Indian Movie Face Database (IMFDB)* created by Shankar Setty et.al. as a benchmark for facial recognition with wide variation. The database consists of thousands of images of 50+ actors taken from more than 100 videos. Since the database has been created manually by cropping the images from the video, there's high variability in terms of pose, expression, illumination, resolution, etc. The original database provides many attributes including:

- Expressions: Anger, Happiness, Sadness, Surprise, Fear, Disgust
- Illumination: Bad, Medium, High
- Pose: Frontal, Left, Right, Up, Down
- Occlusion: Glasses, Beard, Ornaments, Hair, Hand, None, Others
- Age: Child, Young, Middle and Old
- Makeup: Partial makeup, Over-makeup
- Gender: Male, Female

In this scenario, we will use a cleaned and formatted data set with 26742 images split as 19906 train images and 6636 test images respectively. The target here is to use the images and predict the age of the actor/actress within the available classes i.e. young, middle and old making it a multi-class classification problem.

You can download the [train](#) and [test](#) data sets. In each directory, you will find a folder consisting of images along with an excel file which has two columns, ID and Class. The ID column consists of image names like **352.jpg** and Class column holds the respective image character's age like **Old**.

**Data set reference:*

Shankar Setty, Mousa Husain, Parisa Beham, Jyothi Gudavalli, Menaka Kandasamy, Radhesyam Vaddi, Vidyagouri Hemadri, J C Karure, Raja Raju, Rajan, Vijay Kumar and C V Jawahar. "Indian Movie Face Database: A Benchmark for Face Recognition Under Wide Variations"

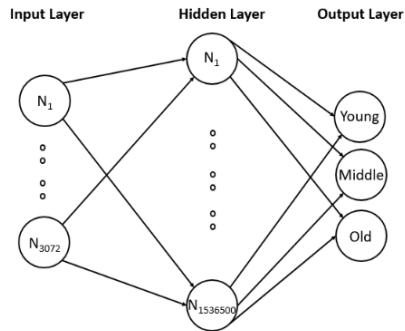
National Conference on Computer Vision, Pattern Recognition, Image Processing and Graphics (NCVPRIPG), 2013.

In the data set, images belong to different sizes and hence cannot be fed directly to the input layer. An input layer has a defined number of nodes which are not changed during the process of implementation of the network. Therefore, before proceeding with building the network, we need to ensure that all the images have equal width and height.

For our problem statement, we will resize all the images to 32 x 32 shape. All the images have red, blue and green color components, therefore, the final shape becomes 32 x 32 x 3 giving us a total of 3072 nodes for the input layer.

Next, we will choose one hidden layer to start with along with 500 nodes making a total of 1536500 (3072 x 500) connections between the input and the hidden layer. We will use the ReLU activation function in this layer.

Next, we have the output layer having only three classes and hence three nodes making a total of 1503 (500 x 3) connections between hidden and output layer. In this layer, we will use the Softmax activation function.



While building the model, we will split the training data into training and validation data set and will find the loss and accuracy for both the data sets.

Now that we have defined the structure of our neural network, let us start implementing the code in Keras.

Before we proceed, let us take a look at the current challenges of the given data set:

- Variations in shape: For example, one image has a shape of (66, 46) whereas another has a shape of (102, 87), there is no consistency
- Multiple viewpoints/ profiles: faces with different viewpoints/profiles may exist
- Brightness and contrast: It varies across images and can introduce discrepancy in few cases
- Quality: Some images are found to be too pixelated

In this resource, we are going to handle the above challenges by performing image preprocessing, as well as implement a basic neural network.

Let us first import all the necessary libraries and modules which will be used throughout the code:

```

1. # Importing necessary Libraries
2. import os
3. import numpy as np
4. import pandas as pd
5. import matplotlib.pyplot as plt
6. %matplotlib inline
7. from sklearn.preprocessing import LabelEncoder
8. from tensorflow.python.keras import utils
9. from keras.models import Sequential
10. from keras.layers import Dense, Flatten, InputLayer
11. import keras
12. import imageio # To read images
13. from PIL import Image # For image resizing

```

Next, let us read the train and test data sets into separate pandas DataFrames as shown below:

```
1. # Reading the data
2. train = pd.read_csv('age_detection_train/train.csv')
3. test = pd.read_csv('age_detection_test/test.csv')
```

Once, both the data sets are read successfully, we can display any random movie character along with their age group to verify the ID against the Class value, as shown below:

```
1. np.random.seed(10)
2. idx = np.random.choice(train.index)
3. img_name = train.ID[idx]
4. img = imageio.imread(os.path.join('age_detection_train/Train', img_name))
5.
6. print('Age group:', train.Class[idx])
7. plt.imshow(img)
8. plt.axis('off')
9. plt.show()
```

Age group: MIDDLE



Next, we can start transforming the data sets to a one-dimensional array after reshaping all the images to a size of 32 x 32 x 3.

Let us reshape and transform the training data first, as shown below:

```
1. temp = []
2. for img_name in train.ID:
3.     img_path = os.path.join('age_detection_train/Train', img_name)
4.     img = imageio.imread(img_path)
5.     img = np.array(Image.fromarray(img).resize((32, 32))).astype('float32')
6.     temp.append(img)
7.
8. train_x = np.stack(temp)
```

Next, let us reshape and transform the testing data, as shown below:

```
1. temp = []
2. for img_name in test.ID:
3.     img_path = os.path.join('age_detection_test/Test', img_name)
4.     img = imageio.imread(img_path)
5.     img = np.array(Image.fromarray(img).resize((32, 32))).astype('float32')
6.     temp.append(img)
7.
8. test_x = np.stack(temp)
```

Next, let us normalize the values in both the data sets to feed it to the network. To normalize, we can divide each value by 255 as the image values lie in the range of 0-255.

```
1. # Normalizing the images
2. train_x = train_x / 255.
3. test_x = test_x / 255.
```

and label encodes the output classes to numerics:

```
1. # Encoding the categorical variable to numeric
2. lb = LabelEncoder()
3. train_y = lb.fit_transform(train.Class)
4. train_y = utils.np_utils.to_categorical(train_y)
```

Next, let us specify the network parameters to be used, as shown below:

```
1. # Specifying all the parameters we will be using in our network
2. input_num_units = (32, 32, 3)
3. hidden_num_units = 500
4. output_num_units = 3
5.
6. epochs = 5
7. batch_size = 128
```

Next, let us define a network with one input layer, one hidden layer, and one output layer, as shown below:

```
1. model = Sequential([
2.     InputLayer(input_shape=input_num_units),
3.     Flatten(),
4.     Dense(units=hidden_num_units, activation='relu'),
5.     Dense(units=output_num_units, activation='softmax'),
6. ])
```

We can also use summary() method to visualize the connections between each layer, as shown below:

```
1. # Printing model summary  
2. model.summary()
```

Layer (type)	Output Shape	Param #
flatten_1 (Flatten)	(None, 3072)	0
dense_1 (Dense)	(None, 500)	1536500
dense_2 (Dense)	(None, 3)	1503
<hr/>		
Total params: 1,538,003		
Trainable params: 1,538,003		
Non-trainable params: 0		

Next, let us compile our network with SGD optimizer and use accuracy as a metric:

```
1. # Compiling and Training Network  
2. model.compile(optimizer='sgd', loss='categorical_crossentropy', metrics=['accuracy'])
```

Now, let us build the model, using the fit() method:

```
1. model.fit(train_x, train_y, batch_size=batch_size, epochs=epochs, verbose=1)
```

This results in the following log:

```
Epoch 1/5  
19906/19906 [=====] - 14s 717us/step - loss: 0.8961 - acc: 0.5817  
Epoch 2/5  
19906/19906 [=====] - 13s 635us/step - loss: 0.8492 - acc: 0.6008  
Epoch 3/5  
19906/19906 [=====] - 13s 630us/step - loss: 0.8316 - acc: 0.6125  
Epoch 4/5  
19906/19906 [=====] - 13s 674us/step - loss: 0.8170 - acc: 0.6206  
Epoch 5/5  
19906/19906 [=====] - 16s 820us/step - loss: 0.8086 - acc: 0.6278
```

We can observe in the above results, that the final accuracy is 62.78%. However, it is recommended that we use 20% to 30% of our training data as a validation data set to observe how the model works on unseen data.

The following code considers 20 percent of the training data as validation data set:

This results in the following log:

```
Train on 15924 samples, validate on 3982 samples  
Epoch 1/5  
15924/15924 [=====] - 11s 669us/step - loss: 0.8014 - acc: 0.6301 - val_loss: 0.7952 - val_acc: 0.6369  
Epoch 2/5  
15924/15924 [=====] - 11s 670us/step - loss: 0.7949 - acc: 0.6355 - val_loss: 0.7872 - val_acc: 0.6477  
Epoch 3/5  
15924/15924 [=====] - 11s 681us/step - loss: 0.7920 - acc: 0.6343 - val_loss: 0.7879 - val_acc: 0.6464  
Epoch 4/5  
15924/15924 [=====] - 11s 709us/step - loss: 0.7867 - acc: 0.6411 - val_loss: 0.7874 - val_acc: 0.6484  
Epoch 5/5  
15924/15924 [=====] - 12s 765us/step - loss: 0.7824 - acc: 0.6451 - val_loss: 0.7965 - val_acc: 0.6364
```

We can observe, that the training accuracy is 64.51% and validation accuracy is 63.64%. Since both the results are quite close we can conclude that there's no overfitting in the model. However, the accuracy itself is too low. The accuracy can be increased by overcoming the previously stated challenges and some difference can even be observed by tuning the hyper-parameters which we are going to observe in the next resource.

With our baseline neural network, we can now predict the age group of test data and save the results in an output file, as shown below:

```
1. # Predicting and importing the result in a csv file
2. pred = model.predict_classes(test_x)
3. pred = lb.inverse_transform(pred)
4.
5. test['Class'] = pred
6. test.to_csv('out.csv', index=False)
```

We can also perform the visual inspection on any random image, as shown below:

```
1. # Visual Inspection of predictions
2. idx = 2481
3. img_name = test.ID[idx]
4.
5. img = imageio.imread(os.path.join('age_detection_test/Test', img_name))
6. plt.imshow(np.array(Image.fromarray(img).resize((128, 128))))
7. pred = model.predict_classes(test_x)
8. print('Original:', train.Class[idx], 'Predicted:', lb.inverse_transform(pred[idx]))
```

Original: MIDDLE Predicted: YOUNG



The network misidentified the current image from the middle age group as young. This could be due to the 64% accuracy of the model. Therefore, let us learn about hyper-parameter tuning and try to improve the results.

11-8-23

Exercise 3: Module name : Understanding and Using CNN : Image recognition Exercise: Design a CNN for Image Recognition which includes hyperparameter tuning.

https://infyspringboard.onwingspan.com/web/en/viewer/web-module/lex_auth_012785694443167744910_shared?collectionId=lex_auth_01274814254931148859_shared&collectionType=Course

Sol:

In this module, we will use a small dataset, CIFAR-10 which can be downloaded from [here](https://infyspringboard.onwingspan.com/common-content-store/Shared/Shared/Public/lex_auth_012782825259556864334_shared/web-hosted/assets/cifar10.zip) along with its [labels](https://infyspringboard.onwingspan.com/common-content-store/Shared/Shared/Public/lex_auth_012782825259556864334_shared/web-hosted/assets/cifar10Labels.csv). This dataset consists of 60,000 images of shape 32x32x3 each categorized into one of the 10 categories (airplanes, cars, birds, cats, deer, dogs, frogs, horses, ships, and trucks). Further, each class has 6000 images.

Few of the images from the dataset are shown below:

Problem Statement

Before we proceed with the learning of a new architecture of the neural network, try to solve the multi-class classification problem using the neural network architecture we used in the previous module on the given CIFAR-10 dataset.

Convolution Neural Network is similar to multi-layer perceptron having made up of neurons with learnable parameters computing the loss function in the last layer. However, CNN architecture primarily made an explicit assumption that inputs are images but in recent times their applications are seen in the areas of text, speech and time series forecasting.

Let's understand the performance difference between the two using a small example. Consider an image with a dimension of 16x16x3.

Implementation using Multi-Layer Perceptron

For a classical neural network, all joined connections in its first hidden layer will result in 768 weights ($16 \times 16 \times 3$). However, with an increased number of neurons and image size (say 300x300x3 with 270000 weights) the structure with a vast number of parameters quickly leads to overfitting. Also, MLP are not known to preserve the spatial features within the images.

Implementation using Convolutional Neural Network

ConvNet pre-assumes the input to be images and hence arranges its neurons in a 3D volume structure: width, height, and depth.

So, an image of shape $16 \times 16 \times 3$ will form a volume of similar shape i.e. $16 \times 16 \times 3$ with neurons connecting only to a slice of preceding layer neurons. The resulting output has a dimension of $1 \times 1 \times h$ where h represents the target labels.

Let us learn what are the components of a CNN and how does it work.

A basic CNN architecture consists of the following layers:

1. Input -> Convolutional -> Pooling -> Output

Among the Convolutional and Pooling layers, both can be repeated as many times as you like.

Input layer

Input layer having two dimensions works as a storage unit for holding raw image data with the preferable size in a multiple of 16, 32, 64, 224, 256, etc. for both height and width for the efficient use of memory fields.

Convolutional Layer

Once the image is loaded in the input layer, the succeeding hidden layers connect back to their preceding layers only on a local region known as the receptive field. This follows a convolution operation which is a combined integration between two functions. It depicts how one function modifies the shape of others.

Since images are represented as a form of a multi-dimensional matrix in the system, therefore, consider the below picture to learn how convolution takes place on a channel (RGB) of an image:

The diagram illustrates the convolution operation from an input image slice to a feature map using a filter.

Input image slice

0	1	1	1	0	0	0
0	0	1	1	1	0	0
0	0	0	1	1	1	0
0	0	0	1	1	0	0
0	0	1	1	0	0	0
0	1	1	0	0	0	0
1	1	0	0	0	0	0

Filter

1	0	1
0	1	0
1	0	1

Feature map

1	4	3	4	1
1	2	4	3	3
1	2	3	4	1
1	3	3	1	1
3	3	1	1	0

* =

Here, at a time a certain image slice is chosen. The filter slides over the input volume convolving with the local region at a time. The number of pixels to jump for next convolution is governed by the stride. Stride with value 1 makes the filter slide over the input volume with 1 pixel, a value of 2 makes the filter slide with 2 pixels and so on. Larger the stride, lesser the spatial extent of output volume. In this illustration, the stride is taken as 1.

Sometimes, the filter size along with stride value doesn't fit the shape of the image, therefore, in such cases, extra padding of zero is preferred. Zero-padding across the input volume border provides us two benefits: first, it helps retain the border information of the image. Since with each convolution, the size of the image keeps reducing and hence without padding the border information may simply be removed. Second, it helps keep the shape of input and output volume equal. Since filter convolution may change the output volume spatial extent, hence padding helps to avoid such cases.

During the process, you can choose the number of filters where each one of them locates distinct features like edges, blobs, etc. Distinct filters are indeed necessary to gain distinct features in the process. For instance, with distinct filters, we can attain features including a sharp image, blurred image, image edges, etc.

To wrap up this idea in one single example, consider the given image where we choose a stride value of two, zero-padding value as one along with two filters. So, to arrive at the output -3 (colored in red), you need to get the sum of the pointwise multiplication of the similar colored matrixes.

For instance,

$$\text{Output volume}_{\text{Red}} = \text{Input volume}_{\text{Green}} * \text{Filter 1}_{\text{Green}} + \text{Input volume}_{\text{Orange}} * \text{Filter 1}_{\text{Orange}} + \text{Input volume}_{\text{Pink}} * \text{Filter 1}_{\text{Pink}} + \text{Bias 1}$$

Similarly, you can proceed to find the values of other cells of the output matrixes. Note, the first output volume matrix is formed using Filter 1 and Bias 1 whereas the second output volume matrix is formed by Filter 2 and Bias 2.

Input volume 7x7x3 (+1 pad)	Filter 1 3x3x3	Filter 2 3x3x3	Output volume 3x3x2
0 0 0 0 0 0 0 0 1 1 0 1 0 0 0 2 0 2 2 2 0 0 1 2 0 1 1 0 0 0 0 0 2 1 0 0 0 1 0 1 1 0 0 0 0 0 0 0 0	1 0 -1 0 -1 -1 -1 -1 0	-1 -1 0 1 0 0 1 1 -1	3 0 -2 -4 -7 -2 -2 -5 0
0 0 0 0 0 0 0 0 0 1 2 0 1 0 0 0 2 2 2 0 0 0 1 2 2 1 1 0 0 0 0 0 1 1 0 0 2 2 0 1 1 0 0 0 0 0 0 0 0	0 0 0 -1 1 0	-1 -1 0 -1 1 0 0 -1 -1	-2 -2 3 -1 -9 -6 1 -4 -5
0 0 0 0 0 0 0 0 0 1 2 1 2 0 0 1 0 0 1 2 0 0 1 0 1 0 2 0 0 1 1 2 1 1 0 0 0 1 1 2 1 0 0 0 0 0 0 0 0	0 -1 -1 0 1 0 1 0 0	-1 0 0 0 0 -1 -1 -1 1	
Bias 1		Bias 2	
1		0	

This layer performs downsampling operation along the two dimensions (width and height), hence reducing the number of required parameters and thus reduced computation and a lesser chance of overfitting. It uses the MAX function and requires two hyperparameters the receptive field, and the stride rate. Padding is generally not used with pooling layer. Also, it doesn't introduce any new parameter as it works on a fixed function.



An alternative to pooling layer: Jost Tobias Springenberg et.al. in their paper "Striving for Simplicity: The All Convolutional Net" suggests that using a higher stride once and using only CONV layers can completely remove the need of having a pooling layer in the architecture.

Fully-Connected layer

Neurons in the Fully-Connected layer are connected to all the activations in previous layers as in ordinary neural networks. It uses the softmax activation function for classifying input images into various classes.

This page lists the conventions required in the process:

If you input a volume of size $W_1 * H_1 * D_1$, it requires four hyperparameters:

1. Number of filters, K
2. Receptive field, F
3. The stride, S
4. The amount of zero-padding, P

Which produces a volume of size $W_2 * H_2 * D_2$ where

- $W_2 = (W_1 - F + 2P) / (S + 1)$
- $H_2 = (H_1 - F + 2P) / (S + 1)$
- $D_2 = K$

Let us start by importing basic modules:

```
1. from matplotlib import pyplot as plt
2. %matplotlib inline
3. from sklearn.preprocessing import LabelEncoder
4. import keras
5. import pandas as pd
6. import numpy as np
7. from PIL import Image
8. import os
9. import warnings
10. warnings.filterwarnings('ignore')
```

Next, let us import the label file and view any random image along with its label:

```
1. labels = pd.read_csv('cifar10_Labels.csv', index_col=0)
2.
3. # View an image
4. img_idx = 5
5. print(labels.label[img_idx])
6. Image.open('cifar10/' + str(img_idx) + '.png')
```

automobile



As we can observe the label is correct as per the image. Now, let us split the data into training and test, follow up with its transformation and normalization:

```

1. # Splitting data into Train and Test data
2. from sklearn.model_selection import train_test_split
3. y_train, y_test = train_test_split(labels.label, test_size=0.3, random_state=42)
4. train_idx, test_idx = y_train.index, y_test.index # Storing indexes for later use
5.
6. # Reading images for training
7. temp = []
8. for img_idx in y_train.index:
9.     img_path = os.path.join('cifar10/', str(img_idx) + '.png')
10.    img = np.array(Image.open(img_path)).astype('float32')
11.    temp.append(img)
12. X_train = np.stack(temp)
13.
14. # Reading images for testing
15. temp = []
16. for img_idx in y_test.index:
17.     img_path = os.path.join('cifar10/', str(img_idx) + '.png')
18.     img = np.array(Image.open(img_path)).astype('float32')
19.     temp.append(img)
20. X_test = np.stack(temp)
21.
22. # Normalizing image data
23. X_train = X_train/255.
24. X_test = X_test/255.
25.

```

The next preprocessing step it to label encode the image respective labels:

```

1. # One-hot encoding 10 output classes
2. encode_X = LabelEncoder()
3. encode_X_fit = encode_X.fit_transform(y_train)
4. y_train = keras.utils.to_categorical(encode_X_fit)

```

Now, let us define the CNN network:

```

1. # Defining CNN network
2.
3. num_classes = 10
4. model = keras.models.Sequential([
5.     # Adding first convolutional Layer
6.     keras.layers.Conv2D(filters=32, kernel_size=(3, 3), strides=1, padding='same', activation='relu',
7.                         kernel_regularizer=keras.regularizers.l2(0.001), input_shape=(32, 32, 3), name='Conv_1'),
8.     # Normalizing the parameters from Last layer to speed up the performance (optional)
9.     keras.layers.BatchNormalization(name='BN_1'),
10.    # Adding first pooling Layer
11.    keras.layers.MaxPool2D(pool_size=(2, 2), name='MaxPool_1'),
12.    # Adding second convolutional Layer
13.    keras.layers.Conv2D(filters=64, kernel_size=(3, 3), strides=1, padding='same', activation='relu',
14.                         kernel_regularizer=keras.regularizers.l2(0.001), name='Conv_2'),
15.    keras.layers.BatchNormalization(name='BN_2'),
16.    # Adding second pooling Layer
17.    keras.layers.MaxPool2D(pool_size=(2, 2), name='MaxPool_2'),
18.    # Flattens the input
19.    keras.layers.Flatten(name='Flat'),
20.    # Fully-Connected Layer
21.    keras.layers.Dense(num_classes, activation='softmax', name='pred_layer')
22. ])

```

In the above model, we have used two convolution layers paired with max pool layers finally connecting with the Fully-Connected layer. We kept the "same" padding i.e., the output volume will have the same length as the original input. For no padding, you can choose the 'valid' argument. The stride is chosen as 1, a total number of 32 and 64 filters for each respective convolution layer and lastly keeping the kernel size as 3x3.

L2 regularization has been added to cost function via the convolution layers. Also, there's an addition of a new concept termed Batch Normalization. It is added due to the following reasons:

- Since we normalize the data before passing it to the input layer to increase the performance, therefore, we add this extra layer of normalization to normalize the values at the intermediate steps.
- It doesn't let the activation go higher or lower, therefore, you can use a higher learning rate to check for new feature possibilities.
- It works in complement to the dropout regularization. It has a slight regularization property as it adds some noise to each hidden layer activations and thus helps avoid overfitting.

Given below is the summary of the above network:

```
1. model.summary()
```

Layer (type)	Output Shape	Param #
Conv_1 (Conv2D)	(None, 32, 32, 32)	896
BN_1 (BatchNormalization)	(None, 32, 32, 32)	128
MaxPool_1 (MaxPooling2D)	(None, 16, 16, 32)	0
Conv_2 (Conv2D)	(None, 16, 16, 64)	18496
BN_2 (BatchNormalization)	(None, 16, 16, 64)	256
MaxPool_2 (MaxPooling2D)	(None, 8, 8, 64)	0
Flat (Flatten)	(None, 4096)	0
pred_layer (Dense)	(None, 10)	40970

Total params: 60,746		
Trainable params: 60,554		
Non-trainable params: 192		

Let us now compile and train the model for just five epochs:

```
1. # Compiling the model
2. model.compile(loss='categorical_crossentropy',
3.                 optimizer=keras.optimizers.Adam(),
4.                 metrics=['accuracy'])
5.
6. cpfile = r'CIFAR10_checkpoint.hdf5' # Weights to be stored in HDF5 format
7. cb_checkpoint = keras.callbacks.ModelCheckpoint(cpfile, monitor='val_acc', verbose=1, save_best_only=True, mode='max')
8. epochs = 5
9. model.fit(X_train, y_train, epochs=epochs, validation_split=0.2, callbacks=[cb_checkpoint])
```

```
Train on 28000 samples, validate on 7000 samples
Epoch 1/5
28000/28000 [=====] - 379s 14ms/step - loss: 1.7326 - acc: 0.4736 - val_loss: 1.7605 - val_acc: 0.4581

Epoch 00001: val_acc improved from -inf to 0.45814, saving model to CIFAR10_checkpoint.hdf5
Epoch 2/5
28000/28000 [=====] - 351s 13ms/step - loss: 1.2379 - acc: 0.6019 - val_loss: 1.7944 - val_acc: 0.4793

Epoch 00002: val_acc improved from 0.45814 to 0.47929, saving model to CIFAR10_checkpoint.hdf5
Epoch 3/5
28000/28000 [=====] - 353s 13ms/step - loss: 1.0563 - acc: 0.6564 - val_loss: 1.2411 - val_acc: 0.6099

Epoch 00003: val_acc improved from 0.47929 to 0.60986, saving model to CIFAR10_checkpoint.hdf5
Epoch 4/5
28000/28000 [=====] - 379s 14ms/step - loss: 0.9514 - acc: 0.6893 - val_loss: 1.3454 - val_acc: 0.5811

Epoch 00004: val_acc did not improve
Epoch 5/5
28000/28000 [=====] - 392s 14ms/step - loss: 0.8666 - acc: 0.7163 - val_loss: 1.2015 - val_acc: 0.6327

Epoch 00005: val_acc improved from 0.60986 to 0.63271, saving model to CIFAR10_checkpoint.hdf5
```

In every batch, the validation accuracy and training accuracy differs much showing a sign of overfitting. However, this model is just to provide you a basic instinct of developing a CNN model from scratch. You can further tune its hyperparameters to increase the performance.

Now, with the given model, let us now perform prediction:

```
1. # << DeprecationWarning: The truth value of an empty array is ambiguous >> can arise due to a NumPy version higher than 1.13.3.
2. # The issue will be updated in upcoming version.
3. pred = encode_X.inverse_transform(model.predict_classes(X_test)[:10])
4. act = y_test[:10]
5.
6. res = pd.DataFrame([pred, act]).T
7. res.columns = ['predicted', 'actual']
8. res
```

	predicted	actual
0	truck	horse
1	ship	ship
2	ship	airplane
3	frog	frog
4	automobile	automobile
5	frog	frog
6	ship	ship
7	airplane	airplane
8	frog	frog
9	ship	dog

We can further proceed with train and test accuracy along with the confusion matrix to judge which class the model is predicting better:

```
1. from mlxtend.evaluate import scoring
2.
3. train_acc = scoring(encode_X.inverse_transform(model.predict_classes(X_train)),
4.                      encode_X.inverse_transform([np.argmax(x) for x in y_train]))
5. test_acc = scoring(encode_X.inverse_transform(model.predict_classes(X_test)), y_test)
6.
7. print('Train accuracy: ', np.round(train_acc, 5))
8. print('Test accuracy: ', np.round(test_acc, 5))
```

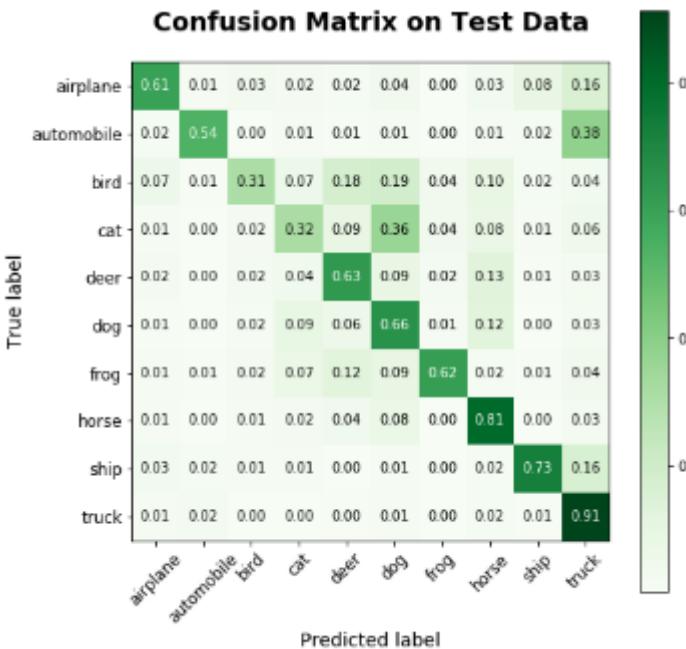
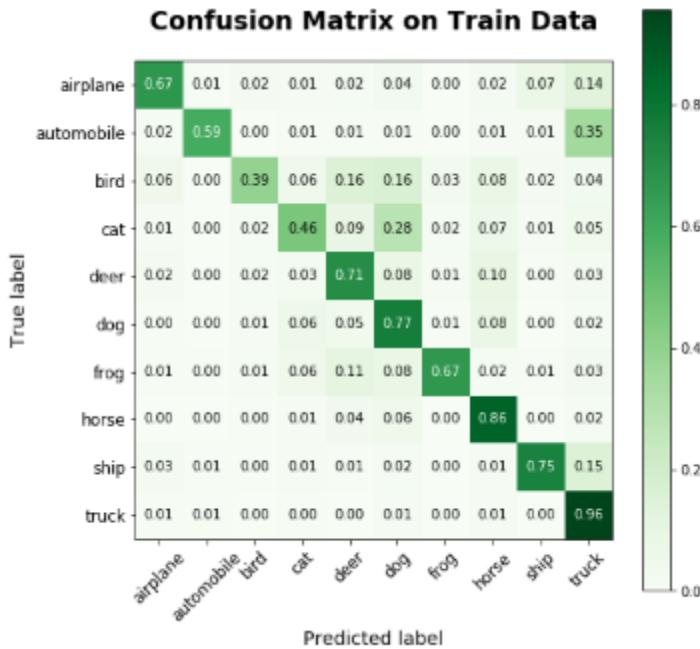
Train accuracy: 0.3176
Test accuracy: 0.3874

```
1. from mlxtend.evaluate import confusion_matrix
2. from mlxtend.plotting import plot_confusion_matrix
3.
4. def plot_cm(cm, text):
5.     class_names=['airplane', 'automobile', 'bird', 'cat', 'deem', 'dog', 'frog', 'horse', 'ship', 'truck']
6.     plot_confusion_matrix(conf_mat=cm,
7.                           colorbar=True, figsize=(8, 8), cmap='Greens',
8.                           show_absolute=False, show_normed=True)
9.     tick_marks = np.arange(len(class_names))
10.    plt.xticks(tick_marks, class_names, rotation=45, fontsize=12)
11.    plt.yticks(tick_marks, class_names, fontsize=12)
12.    plt.xlabel('Predicted label', fontsize=14)
13.    plt.ylabel('True label', fontsize=14)
14.    plt.title(text, fontsize=19, weight='bold')
15.    plt.show()
16.
17. # Train Accuracy
18. train_cm = confusion_matrix(y_target=encode_X.inverse_transform([np.argmax(x) for x in y_train]),
19.                               y_predicted=encode_X.inverse_transform(model.predict_classes(X_train)),
20.                               binary=False)
21. plot_cm(train_cm, 'Confusion Matrix on Train Data')
```

```

23. # Test Accuracy
24. test_cm = confusion_matrix(y_target=y_test,
25.                             y_predicted=encoder_X.inverse_transform(model.predict_classes(X_test)),
26.                             binary=False)
27. plot_cm(test_cm, 'Confusion Matrix on Test Data')

```



The given model has quite a low accuracy on both train and test data, yet its prediction on categories like a truck, horse, and ship is quite remarkable

Exercise 4:

Predicting Sequential Data

Implement a Recurrence Neural Network for Predicting Sequential Data.

Forecasting is the process of predicting the future using current and previous data. The major challenge is understanding the patterns in the sequence of data and then using this pattern to analyse the future. If we were to hand-code the patterns, it would be tedious and changes for the next data. Deep Learning has proven to be better in understanding the patterns in both structured and unstructured data.

To understand the patterns in a long sequence of data, we need networks to analyse patterns across time. Recurrent Networks is the one usually used for learning such data. They are capable of understanding long and short term dependencies or temporal differences.

This post will show you how to implement a forecasting model using LSTM networks in **Keras** and with some cool visualizations. We'll be using the stock price of Google from [yahoo finance](#) but feel free to use any stock data that you like.

Implementation

Use [colab](#) to implement this code to make the visualizations easier, you can use your preferred method. We'll start off with importing necessary libraries:

```
1 import pandas as pd
2 import numpy as np
3 import keras
4 import tensorflow as tf
5 from keras.preprocessing.sequence import TimeseriesGenerator
```

After you've downloaded your .csv file from yahoo finance or your source, load the data using pandas.

```
1 filename = "GOOG.csv"
2 df = pd.read_csv(filename)
3 print(df.info())
```

The info of data frame shows somewhat like this:

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 3797 entries, 0 to 3796
Data columns (total 7 columns):
Date          3797 non-null object
Open           3797 non-null float64
High           3797 non-null float64
Low            3797 non-null float64
Close          3797 non-null float64
Adj Close      3797 non-null float64
Volume         3797 non-null int64
dtypes: float64(5), int64(1), object(1)
memory usage: 207.7+ KB
```

For this tutorial, we require only Date and Close columns, everything else can be dropped.

```
1 df['Date'] = pd.to_datetime(df['Date'])
2 df.set_axis(df['Date'], inplace=True)
3 df.drop(columns=['Open', 'High', 'Low', 'Volume'], inplace=True)
```

Before we do the training and predictions, let's see how the data looks like. For all the visualizations, I'm using the [Plotly](#) python library.

Plotly.... cause its simply the best graphing library and it can produce some good looking graphs.

With plotly, we can define a trace and the layout and it does everything else.



The graph is oscillating from 2018 and the sequence is not smooth... Moving on.

Data Preprocessing

For our analysis, let train the model on the first 80% of data and test it on the remaining 20%.

```

1 close_data = df['Close'].values
2 close_data = close_data.reshape((-1,1))
3
4 split_percent = 0.80
5 split = int(split_percent*len(close_data))
6
7 close_train = close_data[:split]
8 close_test = close_data[split:]
9
10 date_train = df['Date'][:split]
11 date_test = df['Date'][split:]
12
13 print(len(close_train))
14 print(len(close_test))

```

Before we do the training, we need to do some major modification to our data. Remember, our data is still a sequence.. a list of numbers. The neural network is trained as a supervised model. Thus we need to convert the data from sequence to supervised data

Training a neural network of any machine learning model requires the data to be in {<features>,<target>} format. Similarly, we need to convert the given data into this format. Here, we introduce a concept of a look back.

Look back is nothing but the number of previous days' data to use, to predict the value for the next day. For example, let us say look back is 2; so in order to predict the stock price for tomorrow, we need the stock price of today and yesterday.

Coming back to the format, at a given day $x(t)$, the *features* are the values of $x(t-1), x(t-2), \dots, x(t-n)$ where n is look back.

So if our data is like this,

```
[2, 3, 4, 5, 4, 6, 7, 6, 8, 9]
```

The required data format (n=3) would be this:

```
[2, 3, 4] -> [5]
[3, 4, 5] -> [4]
[4, 5, 4] -> [6]
[5, 4, 6] -> [7]
[4, 6, 7] -> [6]
[6, 7, 6] -> [8]
[7, 6, 8] -> [9]
```

There is a module in Keras that does exactly this: [TimeseriesGenerator](#).

```
look_back = 15

train_generator = TimeseriesGenerator(close_train, close_train, length=look_back, batch_size=20)
test_generator = TimeseriesGenerator(close_test, close_test, length=look_back, batch_size=1)
```

Set `look_back` as 15, but you can play around with that value.

Neural Network

Now that our data is ready, we can move on to creating and training our network.

```
1  from keras.models import Sequential
2  from keras.layers import LSTM, Dense
3
4  model = Sequential()
5  model.add(
6      LSTM(10,
7          activation='relu',
8          input_shape=(look_back,1))
9  )
10 model.add(Dense(1))
11 model.compile(optimizer='adam', loss='mse')
12
13 num_epochs = 25
14 model.fit_generator(train_generator, epochs=num_epochs, verbose=1)
```

A simple architecture of LSTM units trained using Adam optimizer and Mean Squared Loss function for 25 epochs. Note that instead of using `model.fit()`, we use `model.fit_generator()` because we have created a data generator.

To know more about LSTM network, see this

<https://colah.github.io/posts/2015-08-Understanding-LSTMs/>

Prediction:

Now that we have completed training, let us see if the network performed well. We can test the model on testing data and see if the prediction and the actual values overlap.

```
1 prediction = model.predict_generator(test_generator)
2
3 close_train = close_train.reshape((-1))
4 close_test = close_test.reshape((-1))
5 prediction = prediction.reshape((-1))
6
7 trace1 = go.Scatter(
8     x = date_train,
9     y = close_train,
10    mode = 'lines',
11    name = 'Data'
12 )
13 trace2 = go.Scatter(
14     x = date_test,
15     y = prediction,
16     mode = 'lines',
17     name = 'Prediction'
18 )
19 trace3 = go.Scatter(
20     x = date_test,
21     y = close_test,
22     mode='lines',
23     name = 'Ground Truth'
24 )
25 layout = go.Layout(
26     title = "Google Stock",
27     xaxis = {'title' : "Date"},
28     yaxis = {'title' : "Close"}
29 )
30 fig = go.Figure(data=[trace1, trace2, trace3], layout=layout)
31 fig.show()
```

Rather than computing loss between predicted and actual values, we can plot it.



From the graph, we can see that prediction and the actual value(ground truth) somewhat overlap. But if you zoom in, the fit is not perfect. We should expect this because it is inevitable as we are performing prediction.

Forecasting

Our testing shows the model is somewhat good. So we can move on to predicting the future or forecasting.

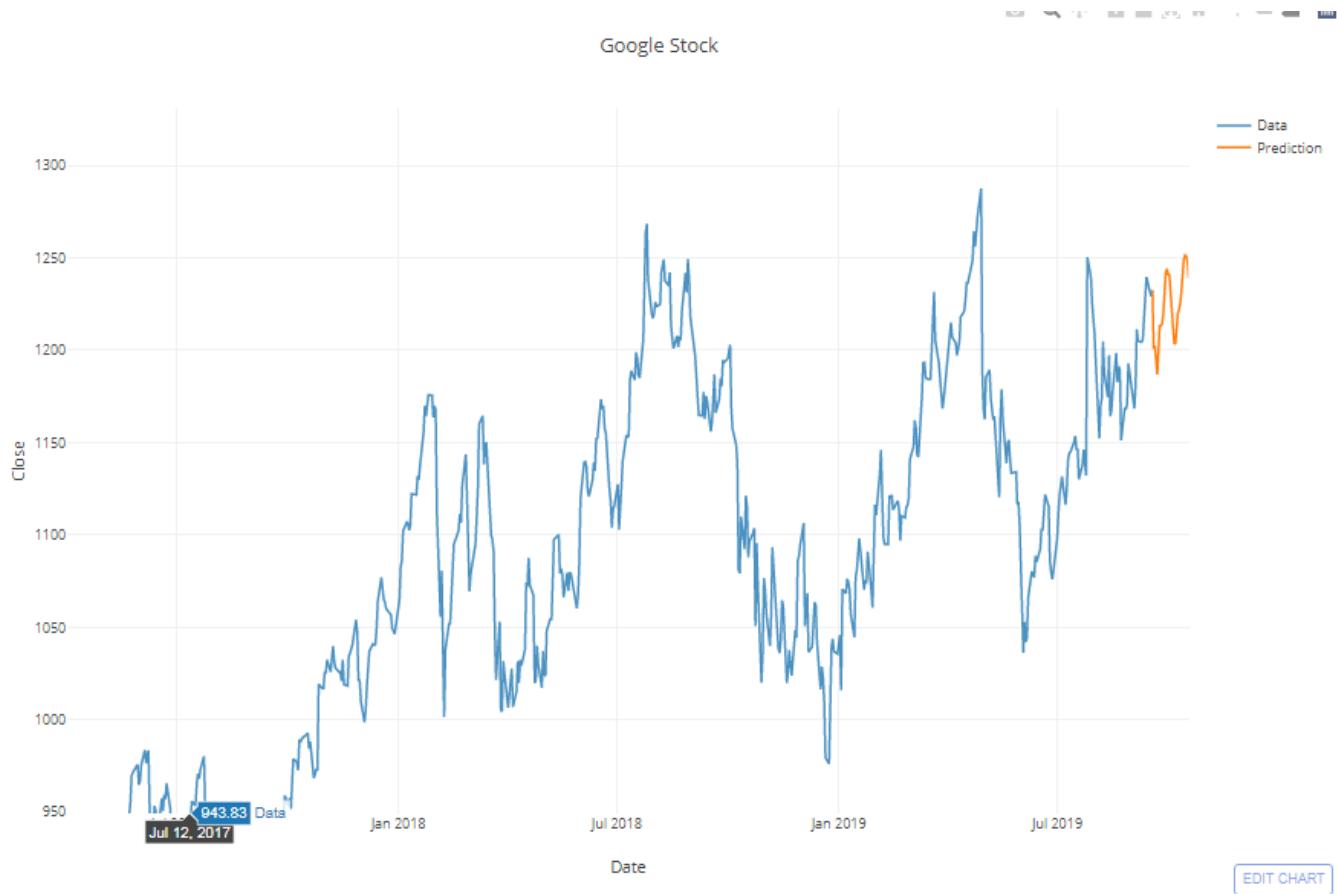
Foreshadowing: Since we are attempting to predict the future, there will be a great amount of uncertainty in the prediction.

Predicting the future is easy... To predict tomorrow's value, feed into the model the past $n(\text{look_back})$ days' values and we get tomorrow's value as output. To get the day after tomorrow's value, feed-in past $n-1$ days' values along with tomorrow's value and the model output day after tomorrow's value.

Forecasting for longer duration is not feasible. So, let's forecast a months stock price.

```
1 close_data = close_data.reshape((-1))
2
3 def predict(num_prediction, model):
4     prediction_list = close_data[-look_back:]
5
6     for _ in range(num_prediction):
7         x = prediction_list[-look_back:]
8         x = x.reshape((1, look_back, 1))
9         out = model.predict(x)[0][0]
10        prediction_list = np.append(prediction_list, out)
11    prediction_list = prediction_list[look_back-1:]
12
13    return prediction_list
14
15 def predict_dates(num_prediction):
16     last_date = df['Date'].values[-1]
17     prediction_dates = pd.date_range(last_date, periods=num_prediction+1).tolist()
18     return prediction_dates
19
20 num_prediction = 30
21 forecast = predict(num_prediction, model)
22 forecast_dates = predict_dates(num_prediction)
```

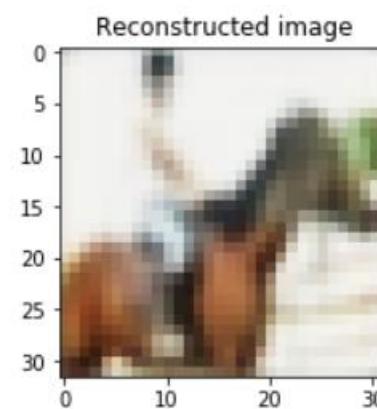
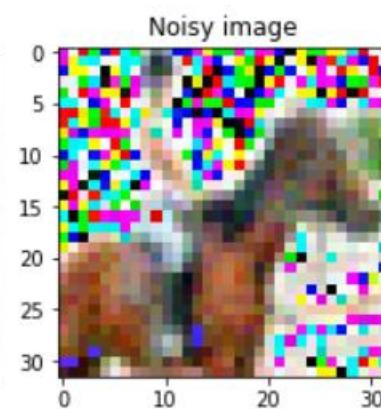
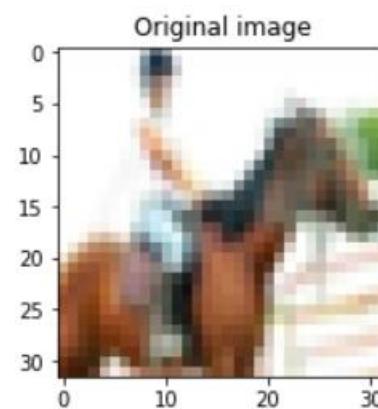
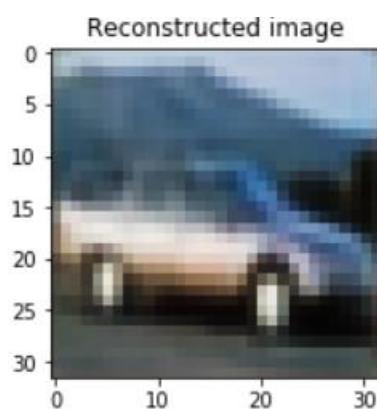
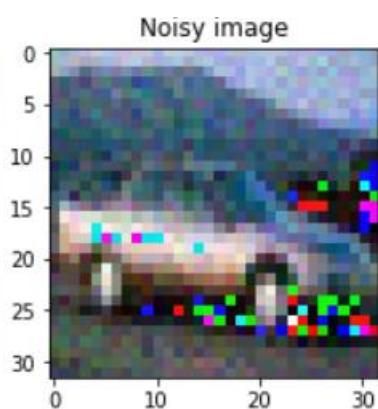
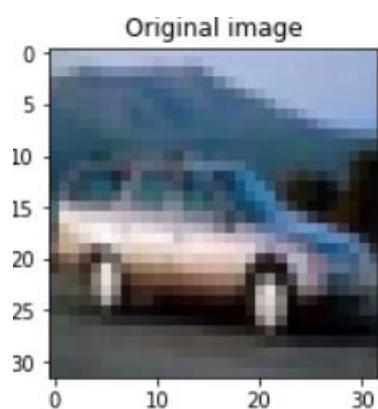
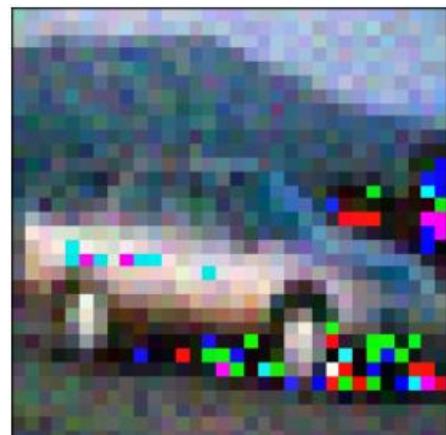
Now plotting the future values,



When predicting the future, there is a good possibility that model output is uncertain to a great extent. The model's output is fed back into it as input. This causes the model's noise and uncertainty to be repeated and amplified.

Conclusion

With this model, we have created a rudimentary model that is able to forecast to a certain extent.



EXERCISE-6

Advanced Deep Learning Architectures

Module Name: Advanced Deep Learning Architectures

Exercise: Implement Object Detection Using YOLO.

OBJECT DETECTION USING YOLO:

- YOLO – “You Only Look Once” are a series of end-to-end deep learning models designed for fast object detection, developed by Joseph Redmon, et al. in the 2015.
- Once the complexity of the image increases, it is not possible to have computational resources to build a Deep Learning model from scratch. So, predefined frameworks and pertained models come in handy. one such framework for object detection is YOLO.
- It's a supremely fast, state-of the art and accurate framework.
- It comes in different versions as shown in fig 5.
- YOLO is implemented on Darknet

How YOLO algorithm works

YOLO architecture is based on CNN and it can be customized according to user's requirement.

Step1: Read the input image



Let, $C =$ number of classes. In the above example, $C = 3$ and the class label are $C1 = \text{Chair}$, $C2 = \text{laptop}$, $C3 = \text{Car}$

Step2: Divide the image into $M \times M$ grid of cells



For each grid cell $X_{ij} \rightarrow Y$, a label Y is calculated. The label Y is a N -dimensional vector, where, N depends on the number of classes. The description of each field is as shown in fig 7. For each grid cell $X_{ij} \rightarrow Y$, a label Y is calculated. The label Y is a N -dimensional vector, where, N depends on the number of classes. The description of each field is as shown in fig 7.

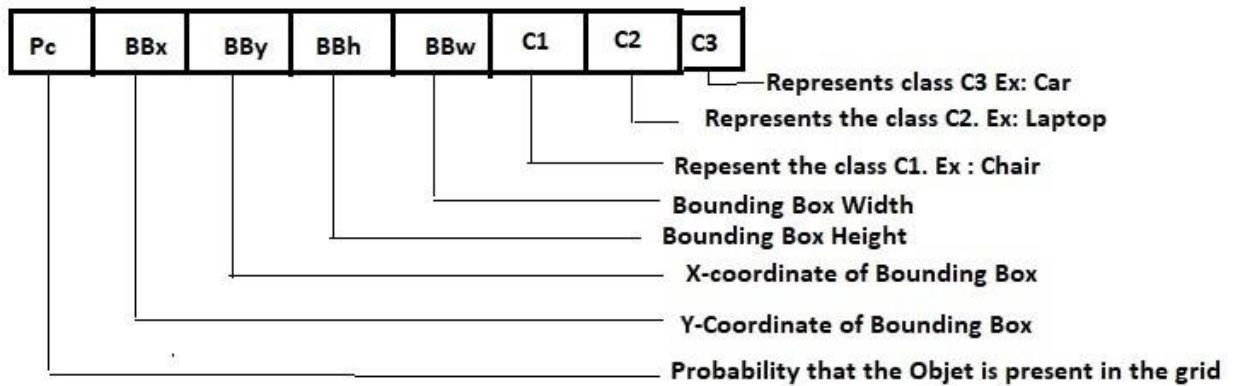


Fig 7. Vector representation of label Y

Step3: Apply Image classification and localization for each grid and predict the bounding box

- The (x, y) coordinates represent the center of the Bounding box relative to the grid cell location and (w, h) – dimension of Bounding box. Both are normalized between [0-1].
- IoU is applied to object detection. Intersection Over Union-IoU is an evaluation metric used to measure the accuracy of an object detector on a dataset.

Step4: Predict the class probabilities of the object

Class probabilities are predicted as P Class Object. This probability is conditioned on the grid cell containing one object.

The vector Y for first grid looks like this

0	BBx	BBy	BBh	BBw	0	0	0
Similarly, the vector Y for grid number 6 look like this:							
1	BBx	BBy	BBh	BBw	1	0	0

The output of this step results in 3x3x8 values i.e., for each grid 8-dimensional vector will be computed.

- In real time scenario the number of grids can be large number like 13x13 and accordingly Y vector varies.

Step5: Train the CNN

The last step is training the Convolutional Neural Network. The normal architecture of CNN is employed with convolutional layer and maxpooling.

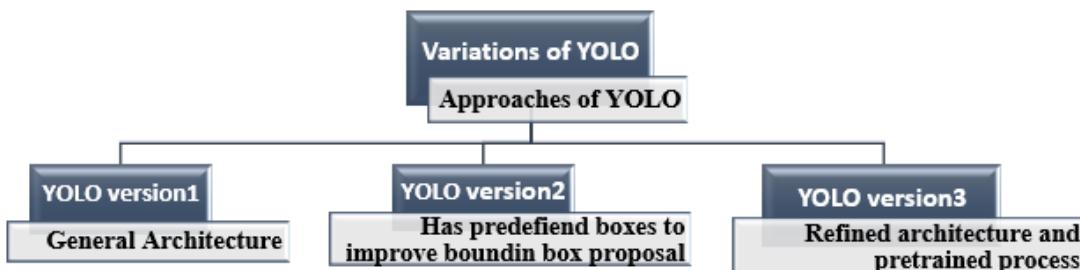


Fig 8. Versions of YOLO

What is Darknet?

- Darknet is an open-source framework that supports Object Detection and Image Classification tasks in the form of Convolutional Neural Networks.
- It is open source and written in C/CUDA
- It is used as the framework for training YOLO, i.e., it sets the architecture of the network
- Darknet is mainly used to implement YOLO algorithm
- The darknet is the executable code.
- This executable code can directly perform object detection in an image, video, camera, and network video stream.

Installation of darknet:

Rule to follow for the successful installation of Darknet:

- Applications should be installed in the correct order for the successful creation of the darknet framework.
- Darknet can be installed with any of the following two optional dependencies namely:

1. In CPU environment using OpenCV (original Darknet Framework, set the GPU flag in Make file when installing darknet to GPU=0.)
2. GPU environment for faster training

1. Steps to install darknet YOLO in CPU execution using OpenCV:



Fig 9. Installation and configuration of environment

1. A clone for the darknet can be created and downloaded from here:
<https://github.com/AlexeyAB/darknet>
2. Extract it to a location of your choice. Darknet take 26.9 MB disk space.
3. Open a MS-PowerShell window in Administrator mode. By executing the command:

<Get-ExecutionPolicy>

```

Windows PowerShell ISE (x86)
File Edit View Tools Debug Add-ons Help
PS C:\Users\meenakshi.h> Get-ExecutionPolicy
RemoteSigned
PS C:\Users\meenakshi.h> |

```

4. If it returns restricted, then run the command below.

PS C:\Users\meenakshi.h> Set-ExecutionPolicy -ExecutionPolicy Unrestricted

```

Windows PowerShell ISE
File Edit View Tools Debug Add-ons Help
Untitled1.ps1 | 1
PS C:\Users\meenakshi.h> set-ExecutionPolicy -ExecutionPolicy Unrestricted

```

- If this command executes correctly, the darknet is installed successfully.

Setting up Pre -Trained models: How to Train YOLO to detect your conventional objects

YOLO v4 Darknet is trained with COCO data set using Convolution Neural Network.

COCO Data set - Common Objects in Context			
No. of classes	Training images	validation images	Download link
80	80,000	40,000	https://cocodataset.org/

Object detection using YOLO is dependent on preparing weights and few configuration files. The weights are pretrained for COCO data set.

Following steps illustrates how to train using YOLO v4:

Download configuration files-yolov4.cfg from here <https://raw.githubusercontent.com/AlexeyAB/darknet/master/cfg/yolov4.cfg> and follow the make few changes in the pretrained parameters.

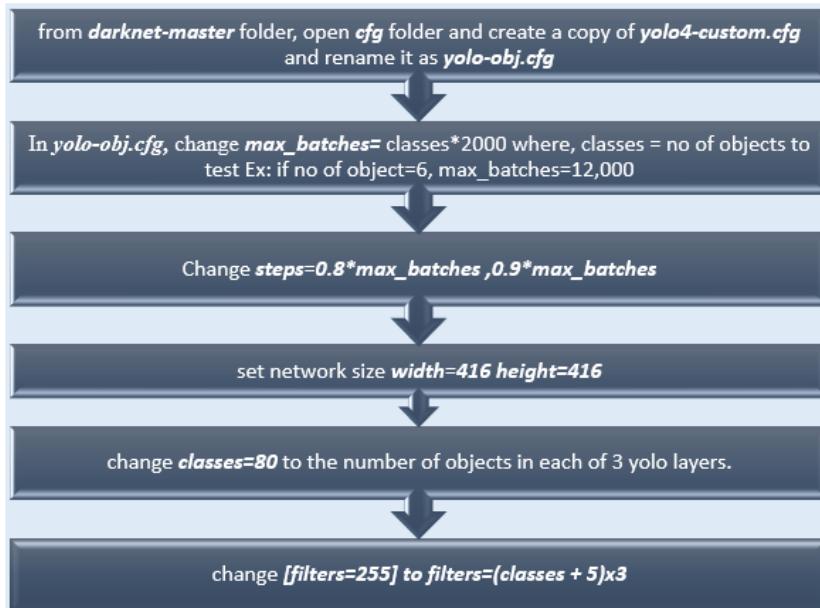


Fig 10. Steps for configuring the darknet files

- Download the pre trained weights from the link [yolov4. Conv .137](#) and save it in the darknet-master folder.
- In a WordPad type the name of each object in separate lines and save the file as obj. names in darknet-master->data folder.
- Create file obj.data in the folder darknet-master->data, and edit the following

```
classes= 3
train = data/train.txt
names = data/obj.names
```

5. Create a folder in darknet-master->data -> obj. Store all the images in obj
6. Create a train.txt file in a path: darknet-master->data folder-> train.txt. This file includes all training images.

data/obj/img1.jpg
data/obj/img2.jpg
data/obj/img3.jpg
data/obj/img4.jpg

7. In the darknet-master folder open Make file in wordpad and change GPU=0, CUDNN=1, OPENCV=1 as shown in the following picture. This is done to make the training on CPU.

Compile darknet:

To compile the darknet execute the following commands:

```
<     make   >
< ./darknet >
```

Train the network:

- The training process could take several hours even days.
- But colab only allow a maximum of 12 hours of running time in ordinary accounts. Those who are interested to train YOLO using darknet in google colab can find the detailshere:
<https://colab.research.google.com/drive/1lTGZsfMaGUpBG4inDIQwIJVW476ibX>
- Training can be done parts by parts. After each 1000 epoch weights are saved in the backup folder so we could just retrain from there. For starting the training run the code.

TESTING: For testing run the following code

```
!./darknet detector test data/obj.data cfg/yolo-obj.cfg backup/yolo-obj_12000.weights
```

What are the Advantages of YOLO over other decoders?

- Rather than using two step method for classification and localization of object, YOLO applies single CNN for both classification and localization of the object.
- YOLO can process images at about 40-90 FPS, so it is quite fast. This means streaming video can be processed in real-time, with negligible latency in a few milliseconds. The

architecture of YOLO makes it extremely fast. When compared with R-CNN, it is 1000 times faster and 100 times faster than fast R-CNN.

Limitations and drawbacks of the YOLO object detector:

1. It especially does not handle objects grouped close together: Since each grid cell predicts only two boxes and can only have one class, this limits the number of nearby objects that YOLO can predict, especially for small objects that appear in groups, such as flocks of birds which may.
2. It does not always handle small objects well: YOLO can detect only 49 objects. The reason for this limitation is due to the YOLO algorithm itself. The YOLO object detector divides an input image into an $M \times M$ grid where each cell in the grid predicts only a single object. If there exist multiple, small objects in a single cell then YOLO will be unable to detect them, ultimately leading to missed object detections.

EXERCISE-7

Optimization of Training in Deep Learning

Module Name: Optimization of Training in Deep Learning

Exercise Name: Design a Deep learning Network for Robust Bi-Tempered Logistic Loss.

Advanced loss functions and parameter tuning:

A **loss function** is used in **neural network** model to optimize the parameter values. Loss function can be classified into two broad categories as shown in fig 2.

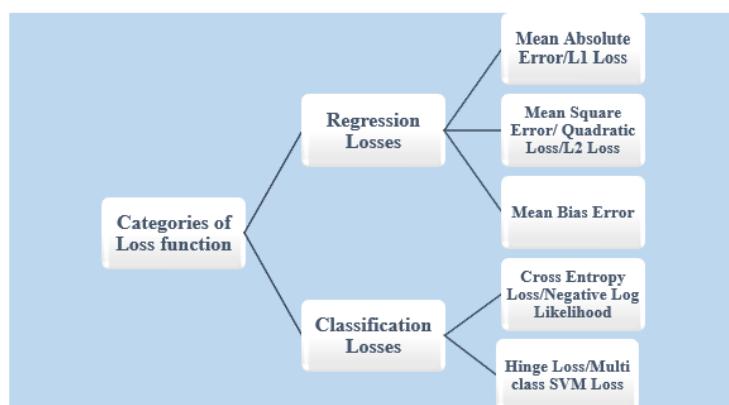


Fig 3. Loss functions

Following table compares the frequently used loss function in deep learning for regression and classification task respectively.

Mean Square Loss (regression)	Cross entropy Loss(classification)
Due to squaring, predictions which are far away from actual values are penalized heavily in comparison to less deviated predictions.	An important aspect of this is that cross entropy loss penalizes heavily the predictions that are <i>confident but wrong</i> .
$MSE = \frac{\sum_{i=1}^n (actual - predicted)^2}{n}$	$\text{cross entropy} = -(actual_i \log(predicted) + (1 - actual_i)\log(1 - actual_i))$
In binary classification model is trained with MSE Cost function, it is not guaranteed to minimize the Cost function.	Distributions with long tails can be modeled poorly with too much weight given to the unlikely events

Demonstration of most used Loss function as optimization algorithm using CNN for MNIST.

Step1: The CNN model is compiled using Adagrad optimizer

```
1. # -*- coding: utf-8 -*-
2. """
```

```

3. Created on Mon Sep 14 01:29:13 2020
4.
5. @author: meenakshi.h
6. """
7.
8. from keras.datasets import mnist
9. import tensorflow
10. from tensorflow.keras.models import Sequential
11. from tensorflow.keras.layers import Dense, Flatten
12. from tensorflow.keras.layers import Conv2D, MaxPooling2D
13. from tensorflow.keras.layers import BatchNormalization

1. # Model configuration
2. batch_size = 250
3. no_epochs = 5
4. no_classes = 10
5. validation_split = 0.2
6. verbosity = 1
7.
8. # Load KMNIST dataset
9. (input_train, target_train), (input_test, target_test) =mnist.load_data()
10.
11. # Shape of the input sets
12. input_train_shape = input_train.shape
13. input_test_shape = input_test.shape

```

```

1. # Keras layer input shape
2. input_shape = (input_train_shape[1], input_train_shape[2], 1)
3.
4. # Reshape the training data to include channels
5. input_train = input_train.reshape(input_train_shape[0],
       input_train_shape[1], input_train_shape[2], 1)
6. input_test = input_test.reshape(input_test_shape[0],
       input_test_shape[1], input_test_shape[2], 1)

1. # Parse numbers as floats
2. input_train = input_train.astype('float32')
3. input_test = input_test.astype('float32')
4.
5. # Normalize input data
6. input_train = input_train / 255
7. input_test = input_test / 255

8.

1. # Create the model
2. model = Sequential()
3. model.add(Conv2D(32, kernel_size=(3, 3), activation='relu',
      input_shape=input_shape))

```

```

4. model.add(BatchNormalization())
5. model.add(MaxPooling2D(pool_size=(2, 2)))
6. model.add(BatchNormalization())
7. model.add(Conv2D(64, kernel_size=(3, 3), activation='relu'))
8. model.add(BatchNormalization())
9.
10. model.add(MaxPooling2D(pool_size=(2, 2)))
11. model.add(BatchNormalization())
12. model.add(Flatten())
13. model.add(Dense(256, activation='relu'))
14. model.add(BatchNormalization())
15. model.add(Dense(no_classes, activation='softmax'))


1.      # Compile the model
2. model.compile(loss=tensorflow.keras.losses.sparse_categorical_crossentropy,
   optimizer=tensorflow.keras.optimizers.Adagrad(
3.      learning_rate=0.001,
4.      initial_accumulator_value=0.1,
5.      epsilon=1e-07))

1. # Fit data to model
2. history = model.fit(input_train, target_train,
3.                       batch_size=batch_size,
4.                       epochs=no_epochs,
5.                       verbose=verbosity,
6.                       validation_split=validation_split)
7.
8. # Generate generalization metric
9. score = model.evaluate(input_test, target_test, verbose=0)
10. print(f'Test loss using Adagrad: {score[0]} / Test accuracy:
    {score[1]}')

```

Step2: Compile the CNN model with Adadelta Optimizer given below replacing Adagrad in the above CNN model

```

1. model.compile(loss=tensorflow.keras.losses.sparse_categorical_crossentropy,
2. optimizer =
   tensorflow.keras.optimizers.Adadelta(learning_rate=0.001,
3. rho=0.95, epsilon=1e-07, name="Adadelta"))

```

Step3: Compile the CNN model with Adam-Adaptive momentum estimation given below replacing Adagrad in the above CNN model

```

1. model.compile(loss=tensorflow.keras.losses.sparse_categorical_crossentropy,
2. optimizer=tensorflow.keras.optimizers.Adam(learning_rate=0.01),
3. metrics=['accuracy'])

```

Step4: Compile the CNN model with Adabound momentum estimation given below replacing Adagrad in the above CNN model

```

1. from keras_adabound import AdaBound
2. model.compile(loss=tensorflow.keras.losses.sparse_categorical_crossentropy,
3. optimizer=AdaBound(lr=1e-3, final_lr=0.1))

```

In each case the model is executed and for first 5 epoch, loss and the accuracy are recorded. In order to study the performance of the optimizer with respect to cross entropy loss function, a graph was plotted. The code for plotting the graph is given below:

```

1. import matplotlib.pyplot as plt
2. x=[1,2,3,4,5]
3.
4. Loss1=[0.6218,0.2478,0.1874,0.158,0.132]
5. accuracy1=[0.8162,0.9332,0.9503,0.9578,0.9599]
6. Loss2=[2.8335,2.2018,1.7276,1.3882,1.1422]
7. accuracy2=[0.1234,0.2605,0.4108,0.5408,0.6399]
8. Loss3=[0.1073,0.0382,0.0261,0.022,0.0158]
9. accuracy3=[0.9673,0.9677,0.9912,0.9926,0.9949]
10. Loss4=[0.1139,0.0281,0.0153,0.009,0.006]
11. accuracy4=[0.9646,0.992,0.9963,0.9984,0.992]
12.
13. fig, axs = plt.subplots(nrows=2, ncols=2, figsize=(6.4, 8.5))
14. #adgrad
15. axs[0, 0].plot(x, Loss1, label="Loss")
16. axs[0, 0].plot(x, accuracy1, label="Accuracy")
17. axs[0, 0].legend()
18. axs[0, 0].grid()
19. axs[0, 0].set(title="Adagrad")
20. #Adadelta
21. axs[0, 1].plot(x, Loss2, label="Loss")
22. axs[0, 1].plot(x, accuracy2, label="Accuracy")
23. axs[0, 1].legend()
24. axs[0, 1].grid()
25. axs[0, 1].set(title="Adadelta")
26. #Adam
27. axs[1, 0].plot(x, Loss3, label="Loss")
28. axs[1, 0].plot(x, accuracy3, label="Accuracy")
29. axs[1, 0].legend()
30. axs[1, 0].grid()
31. axs[1, 0].set(title="Adam")

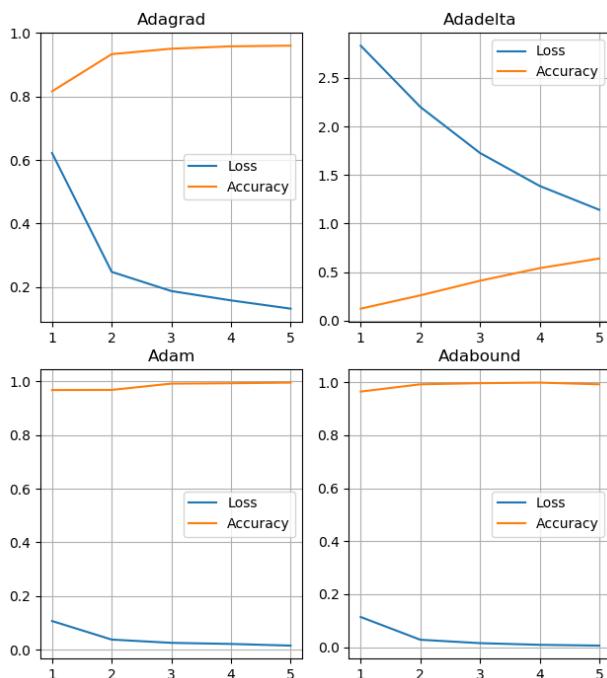
```

```

32. #Adabound
33. axs[1, 1].plot(x, Loss4, label="Loss")
34. axs[1, 1].plot(x, accuracy4, label="Accuracy")
35. axs[1, 1].legend()
36. axs[1, 1].grid()
37. axs[1, 1].set(title="Adabound")
38.
39. fig.tight_layout()
40. plt.show()

```

The below given graph depicts the comparison of Adagrad, Adadelta, Adam and adabound.



When we observe the above graph, we can understand that, Adabound optimization gives the maximum accuracy when as compared to other optimizers and, there is a control in the convergence and generalization.

Loss function can also define as algorithm which converts concepts to practical. So, it is a transformation function on neural networks which converts multiplication of matrix into deep learning. So, far we have seen comprehensive list of loss functions work and let us look at the very recent or advanced loss functions. The advanced loss function in deep learning models are used for specific purposes.

Robust Bi-Tempered Logistic Loss:

We know that, the deep learning model performance is dependent on the quality of training data. The real-world training data sets can be noisy. For example, corrupted images,

mislabeled data are few noisy data sets. The Loss function can fail in handling the noisy training data due to the following two reasons:

1. Highly deviated outliers: Loss function like logistic Loss function is sensitive to outliers

2. Mislabeled data samples: The neural network outputs the class label for each test sample by increasing the distance between the classes. During the process of increasing the decision boundary, the value of the loss function become reduced very fast, so that the training process tend to get close to the boundary of the outliers or mislabeled data samples. Consequently, prediction error occurs. So, a robust loss function is required. “Bi-tempered logistic loss function can be used to generalize he problem of noisy training data.

- As the name says, there are two modifiable parameters that can handle outliers and mislabeled data. They are:
- “temperatures”— t_1 symbolizes the boundedness, and t_2 : indicates the rate of decay in the termination or end of the transfer function
- initialize t_1 and t_2 to 1.0 so that, the logistic loss function is recovered.
- If $t_1 < 1.0$ the boundedness gets increased and if $t_2 > 1.0$ makes transfer function heavy tailed.

How to use Bi-Tempred Logistic Loss:

```
1. #!/bin/bash
2. set -e
3. set -x
4.
5. virtualenv -p python3.
6. source ./bin/activate
7.
8. pip install tensorflow
9. pip install -r bitempered_loss/requirements.txt
10.     python -m bitempered_loss.loss_test
```

Clickhere: <https://ai.googleblog.com/2019/08/bi-tempered-logistic-loss-for-training.html> to know more

The usage of logistic loss using bi-tempered is proved by google for a binary or for two-class classification problem with two-layer on feed-forward neural network.

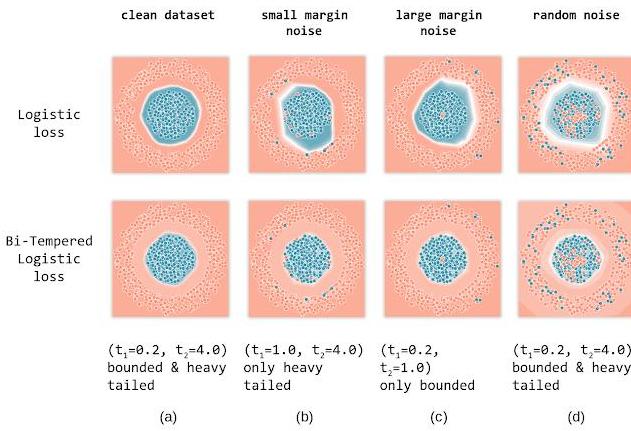


Fig.4.BiTemptedLossfunction [Courtesy: Google AT blog: <https://ai.googleblog.com/2019/08/bi-tempered-logistic-loss-for-training.html>]

GANs Loss Functions:

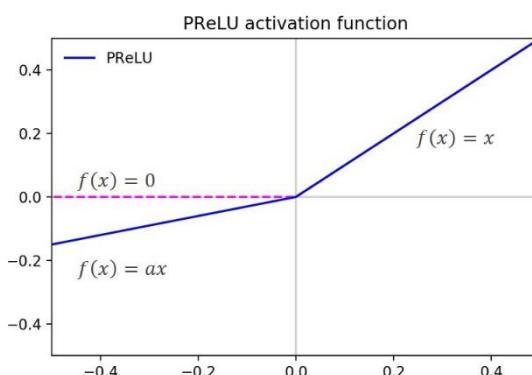
- Generative adversarial networks (GAN) are a powerful subclass of generative models.
- Two main components of GAN are Generator and Discriminator
- This involves the minimization of the generator's loss and maximization of the discriminator's loss.
- Discriminator loss aims at maximizing the probability given to real and fake images. This is a strategy aimed at reducing the worst-case-scenario possible loss. It's simply minimizing the maximum loss. This loss is also used in two-player games to reduce the maximum loss for a layer.

Advanced Activation Functions:

Parametric ReLU: It is a generalization of Leaky ReLU, where, the slope for negative inputs is not predetermined, rather it is considered as a learnable parameter. Formally, it is defined as,

$$y = f(x) = \begin{cases} x, & x > 0 \\ ax, & x \leq 0 \end{cases}$$

where the coefficient 'a' decides the slope of the negative part.



The neural network should learn the best value of itself along with the rest of the model.

Thus, when $a=0$, the function becomes ordinary ReLU. When a is a small and constant value, usually, 0.01, it becomes LReLU.

Even though we are adding one extra parameter to be learnt in this case, it is a negligible overhead when compared to the number of weights that are usually being learnt in deep neural networks.

The gradient of the PReLU function is given as –

$$f'(x) = \frac{\delta f(x)}{\delta a} = \begin{cases} 0, & x > 0 \\ x, & x \leq 0 \end{cases}$$

EXERCISE-8

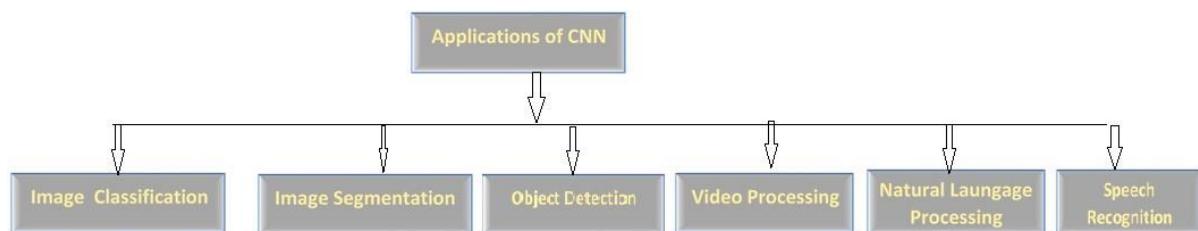
ADVANCED CNN

Module name: Advanced CNN

Exercise: Build AlexNet using Advanced CNN.

INTRODUCTION:

Convolutional Neural Network- CNN is a distinct kind of Neural Networks, which is proved to be the best technique to handle several interesting problems related to Computer Vision and Image Processing.



The learning capability of the deep CNN is mainly because of employing the multiple feature extraction techniques that can automatically extract the patterns from the data. This ability of CNN makes it more powerful and effective CNN.

The voluminous data that is being gathered and the developments in the hardware technology has influenced the new advanced architecture for CNN. Top industries such as Google, AT&T, Amazon, Tesla, Facebook and Microsoft are exploring new architectures of CNN through dynamic research group.

1. ConvNet:

- First multilayered CNN supervised training
- A backpropagation algorithm
- An effective outcome was observed when implemented for handwritten digit recognition and recognition of zip code problems

2. LeNet:

- Improved form of ConvNet
- Used in a document processing for character recognition problems
- Proved to perform optimal in optical character recognition (OCR)and Biometric-fingerprint recognition
- Commercial use in Bank application and ATMs during 1993 and 1996
- Limitation: didn't accomplish better results for other type image recognition problem

3. AlexNet:

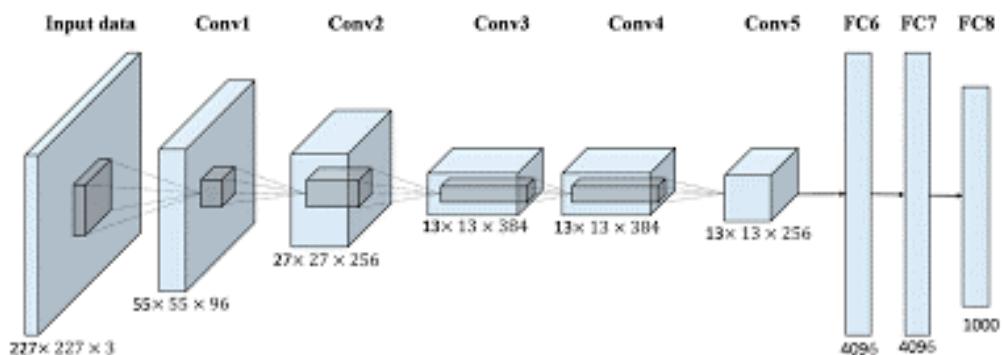
- The main innovation in 2012
- The Idea of basic hyperparameter like stride, filter dimensions, padding, and other was presented for each layer to leverage the image recognition problem

Other like GoogleNet, MobileNet, R-CNN, Fast R-CNN, Faster R-CNN

CASE STUDY OF CNN:

Alex Net:

AlexNet is an incredibly powerful model capable of achieving high accuracies on very challenging datasets. The architecture is as given below figure.



The hyperparameter of AlexNet as listed in below table:

AlexNet architecture Hyperparameter and other details		
Neural Network Layers	Activation Functions	Overfitting Problem
<ul style="list-style-type: none"> • consists of 8 layers • 5 convolutional layers • 3 fully connected layers 	Uses ReLU Nonlinear function instead of tanh function	<p>AlexNet had 60 million parameters, a major issue in terms of overfitting. Two methods were employed to reduce overfitting as given below:</p> <ol style="list-style-type: none"> 1) Data Augmentation: methods like image translations and horizontal reflections, Principle Component Analysis (PCA) on the RGB pixel were used to reduce the error rate 2) Dropout. This technique consists of “turning off” neurons with a predetermined probability. dropout also increases the training time needed for the model’s convergence.

DEMONSTRATION OF ALEXNET :

```

1.      #-----
2. AlexNet Demonstartion
3. #-----
4. #Import keras
5. import numpy as np
6. from keras.datasets import mnist
7. import matplotlib.pyplot as plt
8. #
9. #Load data set
10. (x_train, y_train), (x_test, y_test) = mnist.load_data()
11. print(x_train.shape)

12. print(x_test.shape)

1. element = 200
2. plt.imshow(x_train[element])
3. plt.show()
4. print("Label for the element", element,":", y_train[element])
5. x_train = x_train.reshape((-1, 28*28))
6. x_test = x_test.reshape((-1, 784))
7. print(x_train.shape)
8. print(x_test.shape)
9. x_train = x_train / 255
10. x_test = x_test / 255
11. #

12.

1. from keras.models import Sequential
2. from keras.utils import to_categorical
3. from keras.layers import Dense, Dropout, Activation, Flatten
4. from keras.layers import Conv2D, MaxPooling2D
5. from keras.layers.normalization import BatchNormalization

6. #

```

```

1. # creating model
2. model = Sequential()
3. # 1st Convolutional Layer
4. model.add(Conv2D(filters = 96, input_shape = (60000,784, 3),kernel_size = (11, 11), strides = (4, 4), padding = 'valid'))
5. model.add(Activation('relu'))
6. # Max-Pooling
7. model.add(MaxPooling2D(pool_size = (2, 2),strides = (2, 2), padding = 'valid'))
8. # Batch Normalisation
9. model.add(BatchNormalization())

10.

1. # 2nd Convolutional Layer

```

```
2. model.add(Conv2D(filters = 256, kernel_size = (3, 3),strides = (1, 1), padding = 'valid'))
3. model.add(Activation('relu'))
4. # Max-Pooling
5. model.add(MaxPooling2D(pool_size = (2, 2), strides = (2, 2),padding = 'valid'))
6. # Batch Normalisation
7. model.add(BatchNormalization())

8.

1. # 3rd Convolutional Layer
2. model.add(Conv2D(filters = 384, kernel_size = (3, 3),strides = (1, 1), padding = 'valid'))
3. model.add(Activation('relu'))
4. # Batch Normalisation

5. model.add(BatchNormalization())
```

```
1. # 4th Convolutional Layer
2. model.add(Conv2D(filters = 384, kernel_size = (3, 3),strides = (1, 1), padding = 'valid'))
3. model.add(Activation('relu'))
4. # Batch Normalisation

5. model.add(BatchNormalization())

1. # 5th Convolutional Layer
2. model.add(Conv2D(filters = 256, kernel_size = (3, 3),strides = (1, 1), padding = 'valid'))
3. model.add(Activation('relu'))
4. # Max-Pooling
5. model.add(MaxPooling2D(pool_size = (2, 2), strides = (2, 2), padding = 'valid'))
6. # Batch Normalisation

7. model.add(BatchNormalization())

1. # Flattening

2. model.add(Flatten())
```

```
1. # 1st Dense Layer
2. model.add(Dense(4096, input_shape = (224*224*3, )))
3. model.add(Activation('relu'))
4. # Add Dropout to prevent overfitting
5. model.add(Dropout(0.4))
6. # Batch Normalisation

7. model.add(BatchNormalization())

1. # 2nd Dense Layer
```

```
2. model.add(Dense(4096))
3. model.add(Activation('relu'))
4. # Add Dropout
5. model.add(Dropout(0.4))
6. # Batch Normalisation

7. model.add(BatchNormalization())

1. # Output Softmax Layer
2. model.add(Dense(10))
3. model.add(Activation('softmax'))

4. #-----
```

```
1. # compile the model
2. model.compile(optimizer='Adam',loss='categorical_crossentropy',metrics=['accuracy'])
3. y=to_categorical(y_train)

4. #-----

1. # Fit the model
2. model.fit(x=x_train,y=to_categorical(y_train),epochs=10,batch_size=64,shuffle=True)

3. #-----

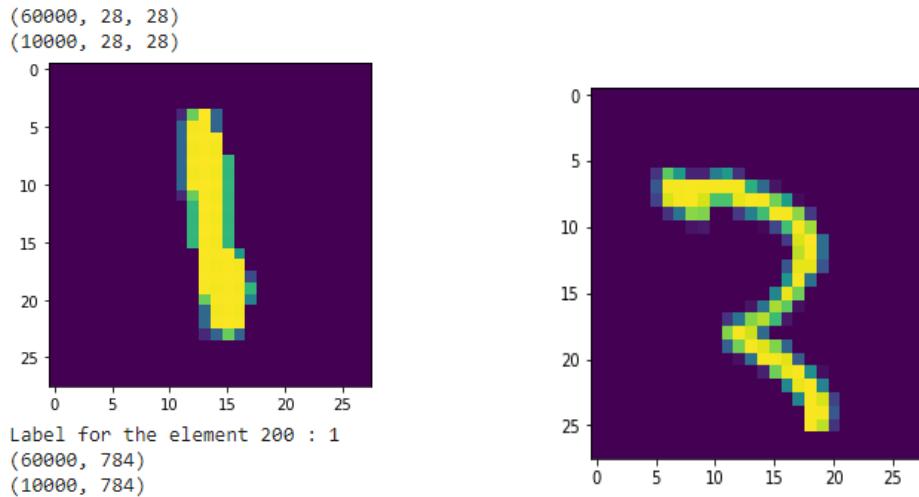
1. # Evaluate the model
2. eval = model.evaluate(x_test, to_categorical(y_test))
3. print('eval')

4. #-----

1. # Predictions
2. predictions = model.predict(x_test[0:100])
3. predictions[0]
4. np.argmax(predictions[0])

5. plt.imshow(x_test[0].reshape(28,28))
```

Output:



```
>>> Instructions for updating:  
>>> Use tf.where in 2.0, which has the same broadcast rule as np.where  
>>> Epoch 1/10  
>>> 60000/60000 [=====] - 21s 344us/step - loss: 0.1972 - acc: 0.9398  
>>> Epoch 2/10  
>>> 60000/60000 [=====] - 20s 332us/step - loss: 0.0802 - acc: 0.9750  
>>> Epoch 3/10  
>>> 60000/60000 [=====] - 19s 312us/step - loss: 0.0530 - acc: 0.9834  
>>> Epoch 4/10  
>>> 60000/60000 [=====] - 19s 311us/step - loss: 0.0390 - acc: 0.9874  
>>> Epoch 5/10  
>>> 60000/60000 [=====] - 19s 314us/step - loss: 0.0300 - acc: 0.9906  
>>> Epoch 6/10  
>>> 60000/60000 [=====] - 19s 319us/step - loss: 0.0254 - acc: 0.9920  
>>> Epoch 7/10  
>>> 60000/60000 [=====] - 19s 309us/step - loss: 0.0218 - acc: 0.9927  
- 19s 309us/step - loss: 0.0218 - acc: 0.9927
```

EXERCISE-9

AUTOENCODERS ADVANCED

Module name: Autoencoders Advanced

Exercise: Demonstration of Application of Autoencoders.

INTRODUCTION:

What are Autoencoders?

Autoencoders are the data encoding techniques based on Unsupervised Artificial Neural Networks. This special type of ANN is trained to encode the data so that in such a way that data is represented in compressed form. The Autoencoders are also trained to decode the data so that, the original data can be reconstructed as far as possible.

Architecture of Autoencoders:

The architecture for autoencoders is varied. In this section LSTM autoencoders is discussed. LSTM based autoencoders are used to encode and decode the sequence data.

Why sequence data is challenging to process?

- Sequence data are challenging for prediction task because the size of the is not fixed but it varies.
- Also, the temporal series of the data representation make it challenging to extract the features.

So, the building a predictive model to predict the sequence data involve sequence of operation and hence such problems are called as Sequence-to Sequence. Autoencoders comes as the best choice to handle sequence-to-sequence problems.

DEMONSTRATION OF APPLICATION OF AUTOENCODERS:

LSTM based autoencoders can be created to for various applications. Some of them are demonstrated below.

1. Reconstruction of sequence using Autoencoders:

Step1: Building a simple autoencoders to create simple sequence

1. from numpy import array
2. from keras. models import Sequential
3. from keras. layers import LSTM
4. from keras. layers import Dense
5. from keras. layers import RepeatVector

```

6. from keras.layers import TimeDistributed
7. from keras.utils import plot_model
8. # lstm autoencoder recreate sequence
9.
10. # define input sequence
11. sequence = array ([0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9])
12. # Reshape input into [samples, timesteps, features]
13. n_in = len(sequence)
14. sequence = sequence.reshape ((1, n_in, 1))
15. # define model
16. model = Sequential ()
17. model.add (LSTM (100, activation='relu', input_shape=(n_in,1)))
18. model.add (RepeatVector(n_in))
19. model.add (LSTM (100, activation='relu', return_sequences=True))
20. model.add (TimeDistributed (Dense(1)))
21. model.compile(optimizer='adam', loss='mse')
22. # fit model
23. model.fit(sequence, sequence, epochs=300, verbose=0)
24. plot_model(model, show_shapes=True, to_file='reconstruct_lstm_autoencoder.png')
25. # demonstrate recreation
26. yhat = model.predict(sequence, verbose=0)

27. print(yhat[0,:,0])

```

2. Prediction of the sequence of number using Autoencoders:

Like reconstruction, autoencoders can be used to predict the sequence, the code is as given below:

```

1. # lstm autoencoder predict sequence
2. from numpy import array
3. from keras.models import Sequential
4. from keras.layers import LSTM
5. from keras.layers import Dense
6. from keras.layers import RepeatVector
7. from keras.layers import TimeDistributed
8. from keras.utils import plot_model
9. # define input sequence
10. seq_in = array([0.1, 0.2, 0.3, 0.4, 0.5, 0.6, 0.7, 0.8, 0.9])
11. # reshape input into [samples, timesteps, features]
12. n_in = len(seq_in)
13. seq_in = seq_in.reshape((1, n_in, 1))
14. # prepare output sequence
15. seq_out = seq_in[:, 1:, :]
16. n_out = n_in - 1
17. # define model
18. model = Sequential()
19. model.add(LSTM(100, activation='relu', input_shape=(n_in,1)))
20. model.add(RepeatVector(n_out))
21. model.add(LSTM(100, activation='relu', return_sequences=True))
22. model.add(TimeDistributed(Dense(1)))

```

```

23. model.compile(optimizer='adam', loss='mse')
24. plot_model(model, show_shapes=True, to_file='predict_lstm_autoencoder.png')
25. #fit model
26. model.fit(seq_in, seq_out, epochs=300, verbose=0)
27. #demonstrate prediction
28. yhat = model.predict(seq_in, verbose=0)

29. print(yhat[0,:,:])

```

3. Outlier/Anomaly detection using Autoencoders:

Suppose the input data is highly correlated and requires a technique to detect the anomaly or an outlier then, Autoencoders is the best choice. Since, autoencoders can encode the data in the compressed form, they can handle the correlated data.

Let's train the autoencoders using MNIST data set using simple Feed Forward neural network.

Code: Simple 6 layered feed forward Autoencoders

```

Downloading data from https://storage.googleapis.com/tensorflow/tf-keras-datasets/mnist.npz
11493376/11490434 [=====] - 0s 0us/step
Epoch 1/10
59/59 [=====] - 7s 114ms/step - loss: 0.0758 - val_loss: 0.0516
Epoch 2/10
59/59 [=====] - 6s 110ms/step - loss: 0.0440 - val_loss: 0.0369
Epoch 3/10
59/59 [=====] - 7s 111ms/step - loss: 0.0342 - val_loss: 0.0311
Epoch 4/10
59/59 [=====] - 7s 111ms/step - loss: 0.0297 - val_loss: 0.0274
Epoch 5/10
59/59 [=====] - 7s 110ms/step - loss: 0.0267 - val_loss: 0.0251
Epoch 6/10
59/59 [=====] - 6s 110ms/step - loss: 0.0246 - val_loss: 0.0233
Epoch 7/10
59/59 [=====] - 6s 110ms/step - loss: 0.0231 - val_loss: 0.0221
Epoch 8/10
59/59 [=====] - 6s 110ms/step - loss: 0.0220 - val_loss: 0.0211
Epoch 9/10
59/59 [=====] - 7s 110ms/step - loss: 0.0211 - val_loss: 0.0204
Epoch 10/10
59/59 [=====] - 7s 113ms/step - loss: 0.0203 - val_loss: 0.0195

```

Once the autoencoders is trained on MNIST data set, an anomaly detection can be done using 2 different images. First one of the images from the MNIST data set is chosen and feed to the trained autoencoders. Since, this image is not an anomaly, the error or loss function is expected to be very low. Next, when some random image is given as test image, the loss rate is expected to be very high as it is an anomaly.

Simple 6 layered Autoencoders build to train on MNIST data

1. import numpy as np
2. import keras
3. from keras.datasets import mnist
4. from keras.models import Sequential, Model
5. from keras.layers import Dense, Input

```

6. from keras import optimizers
7. from keras.optimizers import Adam
8.
9. (x_train, y_train), (x_test, y_test) = mnist.load_data()
10. train_x = x_train.reshape(60000, 784) / 255
11. val_x = x_test.reshape(10000, 784) / 255
12.
13. autoencoder = Sequential()
14. autoencoder.add(Dense(512, activation='elu', input_shape=(784,)))
15. autoencoder.add(Dense(128, activation='elu'))
16. autoencoder.add(Dense(10, activation='linear', name="bottleneck"))
17. autoencoder.add(Dense(128, activation='elu'))
18. autoencoder.add(Dense(512, activation='elu'))
19. autoencoder.add(Dense(784, activation='sigmoid'))
20. autoencoder.compile(loss='mean_squared_error', optimizer=Adam())
21. trained_model = autoencoder.Fit(train_x, train_x, batch_size=1024, epochs=10, verbose=1,
   validation_data=(val_x, val_x))
22. encoder = Model(autoencoder.Input, autoencoder.get_layer('bottleneck').Output)
23. encoded_data = encoder.Predict(train_x) # bottleneck representation
24. decoded_output = autoencoder.Predict(train_x) # reconstruction
25. encoding_dim = 10
26.
27. # Return the decoder
28. encoded_input = Input(shape=(encoding_dim,))
29. decoder = autoencoder.layers[-3](encoded_input)
30. decoder = autoencoder.layers[-2](decoder)
31. decoder = autoencoder.layers[-1](decoder)

32. decoder = Model(encoded_input, decoder)

```

Anomaly Detection

```

1. # %matplotlib inline
2. from keras.preprocessing import image
3. # if the img.png is not one of the MNIST dataset that the model was trained on, the error will be
   very high.
4. img = image.load_img("C:\Users\meenakshi.h\Desktop\Images\fig12.png", target_size=(28, 28),
   color_mode = "grayscale")
5. input_img = image.img_to_array(img)
6. inputs = input_img.reshape(1,784)
7. target_data = autoencoder.predict(inputs)
8. dist = np.linalg.norm(inputs - target_data, axis=-1)

9. print(dist)

```

EXERCISE-10

Advanced GANs

Module name: Advanced GANs

Exercise: Demonstration of GAN

INTRODUCTION:

Can we train the machine to generate the following?

- Given the image of a criminal whose face is masked with glass. Is it possible for the machine to automatically generate the photo without glass?
- Suppose few simple sentences which describe the bird as: "*A small bird with red head. The features fade from red to gray from head to tail.*" are given to machine. Is it possible for the machine to generate the image of the bird?
- Given a photo of a person which is captured 10 years ago, is it possible for the machine to generate a photo of the same person for the present time?

Yes, advanced deep learning has made an amazing progress, where machines can be trained to generate or create new images. Deep learning models which can generate new images from the given training data set are called as **Generative Adversarial Networks** or in short **GANs**. Before understanding GANs it is required to have the knowledge of Generative models.

Difference between Discriminative and Generative: Machine learning models can broadly fall under two categories namely, discriminative and generative. Following table lists the difference between discriminative and generative models.

Discriminative	Generative
Discriminative models are developed to predict a class label using the given labelled training set.	Generative models are the one that are used to create or generate new examples based on the input data distribution
Since, the samples are discriminated based on the labelled data set, they are Supervised Machine learning in nature.	Since, the input data is not labelled, generative models are called as Unsupervised machine learning
Example: Classification, logistic regression	Example: Naive Bayes can be Discriminative and generative, Latent Dirichlet Allocation-LDA, Gaussian Mixture Model-GMM and deep learning based model namely GANs- Generative Adversarial Network

The advanced Deep learning unsupervised Generative models that are the state-of the art include:

- 1) Variational Autoencoder, or VAE, and
- 2) The Generative Adversarial Network, or GAN

Let's look at the architectural aspects of GANs

ARCHITECTURE AND TYPES OF GAN:

The architecture of the deep learning GAN models consists of two modules namely **generator** and **Discriminator**.

1. A generator model:

- It is the learning component of a GAN models.
- It learns to generate the new data by incorporating the feedback received from the discriminator.
- It learns to allow the discriminator to classify its newly generated data as real.
- Hence, training the Generator also requires the discriminator to be considered.

2. A Discriminator models:

- It is a classifier in GANs
- Its job is to classify the output of the generator (newly generated data) from the real one.

The below given figure 1. Shows the architecture of GANs

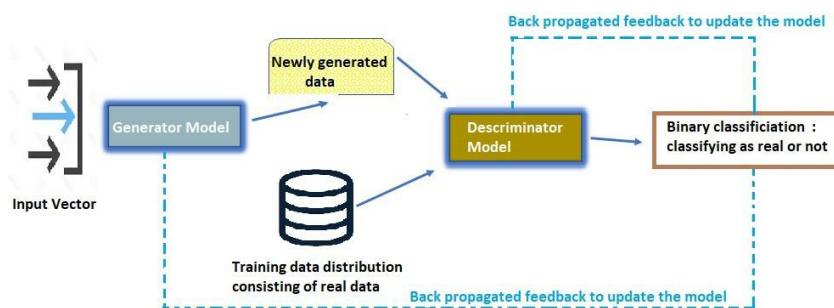


Fig 1. Architecture of GANs

DEMONSTRATION OF GAN:

The application of GAN models is varied. Here two case studies are considered to demonstrate GAN for Data set generation. They are as follows:

1. Image Augmentation using MNIST data set
2. New image generation for CIFAR data Set

1. Image Augmentation: Case study of GAN

Whenever the data set don't have enough samples to train the machine due to various constraints in data collection process then, it becomes necessary to use Augmentation. Particularly, when more complex object needs to be recognized then, Image data augmentation technique is used. It is a method of artificially escalating the size of a training dataset by creating artificially new set of images.

Using Keras Image augmentation is demonstrated by building deep learning GANs.

ImageDataGenerator class available in Keras is used in demonstration. It defines the configuration for image data preparation and augmentation.

In this demonstration following properties are demonstrated:

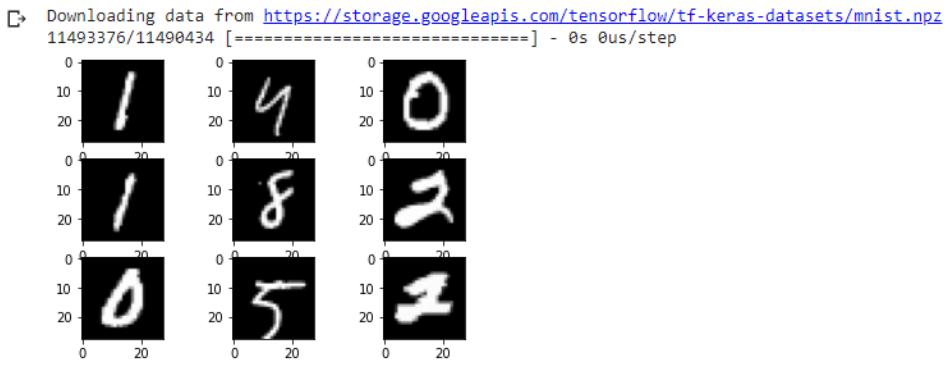
- Feature standardization.
- ZCA whitening.
- Random flips.

a. Feature Standardization

Using GANs model the pixel values across the entire dataset can be standardized. Feature standardization is the process of standardizing the pixel which is performed for each column in a tabular dataset. This can be done by setting the feature wise_center and feature wise_std_normalization arguments on the ImageDataGenerator class.

```
1. from keras.datasets import mnist
2. from keras.preprocessing.image import ImageDataGenerator
3. from matplotlib import pyplot
4. # load data
5. (X_train, y_train), (X_test, y_test) = mnist.load_data()
6. # reshape to be [samples][width][height][channels]
7. X_train = X_train.reshape((X_train.shape[0], 28, 28, 1))
8. X_test = X_test.reshape((X_test.shape[0], 28, 28, 1))
9. # convert from int to float
10. X_train = X_train.astype('float32')
11. X_test = X_test.astype('float32')
12. # define data preparation
13. datagen = ImageDataGenerator(featurewise_center=True,
   featurewise_std_normalization=True)
14. # fit parameters from data
15. datagen.fit(X_train)
16. # configure batch size and retrieve one batch of images
17. for X_batch, y_batch in datagen.flow(X_train, y_train, batch_size=9):
18.     # create a grid of 3x3 images
19.     for i in range(0, 9):
20.         pyplot.subplot(330 + 1 + i)
21.         pyplot.imshow(X_batch[i].reshape(28, 28), cmap=pyplot.get_cmap('gray'))
22.     # show the plot
23.     pyplot.show()
24.     break
```

Output



b. ZCA-Zero Component Analysis Whitening

Suppose the pixel has many redundant pixels then, training process can't be effective. So, to reduce the redundant pixels whitening of an image is used. The process of transforming the original image using a linear algebra operation that reduces the redundancy in the matrix of pixel is called as Whitening transformation.

Advantage of whitening: Less redundant pixels in the image is expected to improve the structures and features in the image so that, machine can learn image effectively.

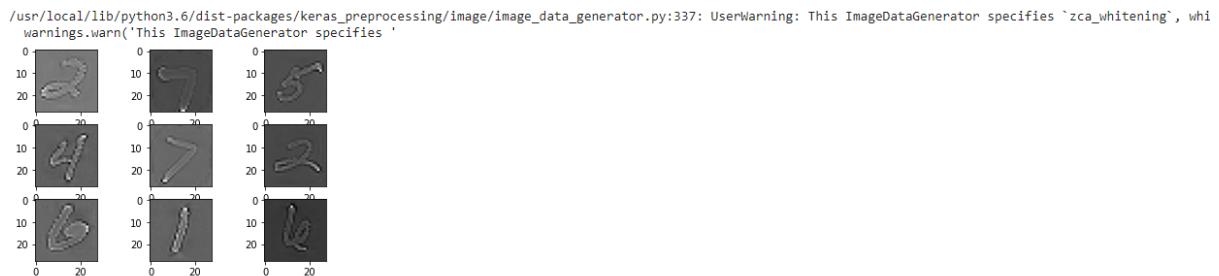
In this demonstration, ZCA is used to show GANs application in generating new image after eliminating the redundant pixels.

```

1. # ZCA whitening
2. from keras.datasets import mnist
3. from keras.preprocessing.image import ImageDataGenerator
4. from matplotlib import pyplot
5. # load data
6. (X_train, y_train), (X_test, y_test) = mnist.load_data()
7. # reshape to be [samples][width][height][channels]
8. X_train = X_train.reshape((X_train.shape[0], 28, 28, 1))
9. X_test = X_test.reshape((X_test.shape[0], 28, 28, 1))
10. # convert from int to float
11. X_train = X_train.astype('float32')
12. X_test = X_test.astype('float32')
13. # define data preparation
14. datagen = ImageDataGenerator(zca_whitening=True)
15. # fit parameters from data
16. datagen.fit(X_train)
17. # configure batch size and retrieve one batch of images
18. for X_batch, y_batch in datagen.flow(X_train, y_train, batch_size=9):
19.     # create a grid of 3x3 images
20.     for i in range(0, 9):
21.         pyplot.subplot(330 + 1 + i)
22.         pyplot.imshow(X_batch[i].reshape(28, 28), cmap=pyplot.get_cmap('gray'))
23.     # show the plot
24.     pyplot.show()
25.     break

```

Output :

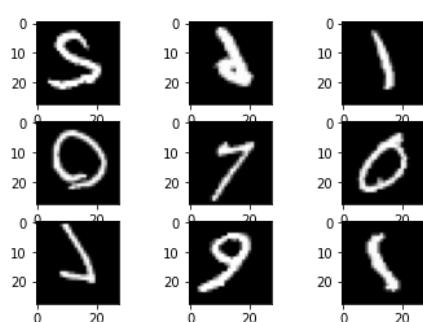


c. Random Flips

Random Flip can be used as augmentation technique on an image data to improve the performance on large and complex problems.

```
1. # Random Flips
2. from keras.datasets import mnist
3. from keras.preprocessing.image import ImageDataGenerator
4. from matplotlib import pyplot
5. # load data
6. (X_train, y_train), (X_test, y_test) = mnist.load_data()
7. # reshape to be [samples][width][height][channels]
8. X_train = X_train.reshape((X_train.shape[0], 28, 28, 1))
9. X_test = X_test.reshape((X_test.shape[0], 28, 28, 1))
10. # convert from int to float
11. X_train = X_train.astype('float32')
12. X_test = X_test.astype('float32')
13. # define data preparation
14. datagen = ImageDataGenerator(horizontal_flip=True, vertical_flip=True)
15. # fit parameters from data
16. datagen.fit(X_train)
17. # configure batch size and retrieve one batch of images
18. for X_batch, y_batch in datagen.flow(X_train, y_train, batch_size=9):
19.     # create a grid of 3x3 images
20.     for i in range(0, 9):
21.         pyplot.subplot(330 + 1 + i)
22.         pyplot.imshow(X_batch[i].reshape(28, 28), cmap=pyplot.get_cmap('gray'))
23.     # show the plot
24.     pyplot.show()
25.     break
```

Output:



New image generation using CIFAR data Set

Step 1: Importing the required libraries

```
1. import numpy as np
2. import matplotlib.pyplot as plt
3. import keras
4. from keras.layers import Input, Dense, Reshape, Flatten, Dropout
5. from keras.layers import BatchNormalization, Activation, ZeroPadding2D
6. from keras.layers.advanced_activations import LeakyReLU
7. from keras.layers.convolutional import UpSampling2D, Conv2D
8. from keras.models import Sequential, Model

9. from keras.optimizers import Adam,SGD
```

Step 2: Loading the data

```
1. #Loading the CIFAR10 data
2. (X, y), (_, _) = keras.datasets.cifar10.load_data()
3. #Selecting a single class images
4. #The number was randomly chosen and any number
5. #between 1 to 10 can be chosen

6. X = X[y.flatten() == 8]
```

Step 3: Defining parameters to be used in later processes

```
1. #Defining the Input shape
2. image_shape = (32, 32, 3)

3. latent_dimensions = 100
```

Step 4: Defining a utility function to build the Generator

```
1. def build_generator():
2.     model = Sequential()
3.     #Building the input layer
4.     model.add(Dense(128 * 8 * 8, activation="relu",
5.                     input_dim=latent_dimensions))
6.     model.add(Reshape((8, 8, 128)))
7.     model.add(UpSampling2D())
8.     model.add(Conv2D(128, kernel_size=3, padding="same"))
9.     model.add(BatchNormalization(momentum=0.78))
10.    model.add(Activation("relu"))
11.    model.add(UpSampling2D())
12.    model.add(Conv2D(64, kernel_size=3, padding="same"))
13.    model.add(BatchNormalization(momentum=0.78))
```

```

14.     model.add(Activation("relu"))
15.     model.add(Conv2D(3, kernel_size=3, padding="same"))
16.     model.add(Activation("tanh"))
17. #Generating the output image
18.     noise = Input(shape=(latent_dimensions,))
19.     image = model(noise)

20. return Model(noise, image)

```

Step 5: Defining a utility function to build the Discriminator

```

1. def build_discriminator():
2.     #Building the convolutional layers
3.     #to classify whether an image is real or fake
4.     model = Sequential()
5.     model.add(Conv2D(32, kernel_size=3, strides=2,
6.                     input_shape=image_shape, padding="same"))
7.     model.add(LeakyReLU(alpha=0.2))
8.     model.add(Dropout(0.25))
9.     model.add(Conv2D(64, kernel_size=3, strides=2, padding="same"))
10.    model.add(ZeroPadding2D(padding=((0,1),(0,1))))
11.    model.add(BatchNormalization(momentum=0.82))
12.    model.add(LeakyReLU(alpha=0.25))
13.    model.add(Dropout(0.25))
14.    model.add(Conv2D(128, kernel_size=3, strides=2, padding="same"))
15.    model.add(BatchNormalization(momentum=0.82))
16.    model.add(LeakyReLU(alpha=0.2))
17.    model.add(Dropout(0.25))
18.    model.add(Conv2D(256, kernel_size=3, strides=1, padding="same"))
19.    model.add(BatchNormalization(momentum=0.8))
20.    model.add(LeakyReLU(alpha=0.25))
21.    model.add(Dropout(0.25))
22.    #Building the output layer
23.    model.add(Flatten())
24.    model.add(Dense(1, activation='sigmoid'))
25.    image = Input(shape=image_shape)
26.    validity = model(image)

27. return Model(image, validity)

```

Step 6: Defining a utility function to display the generated images

```

1. def display_images():
2.     r, c = 4,4
3.     noise = np.random.normal(0, 1, (r * c,latent_dimensions))
4.     generated_images = generator.predict(noise)
5. #Scaling the generated images
6.     generated_images = 0.5 * generated_images + 0.5
7.     fig, axs = plt.subplots(r, c)

```

```

8.     count = 0
9.     for i in range(r):
10.        for j in range(c):
11.            axs[i,j].imshow(generated_images[count, :, :, :])
12.            axs[i,j].axis('off')
13.            count += 1
14.    plt.show()

15.    plt.close()

```

Step 7: Building the Generative Adversarial Network

```

1. # Building and compiling the discriminator
2. discriminator = build_discriminator()
3. discriminator.compile(loss='binary_crossentropy',
4.                         optimizer=Adam(0.0002,0.5),
5.                         metrics=['accuracy'])
6.
7. #Making the Discriminator untrainable
8. #so that the generator can learn from fixed gradient
9. discriminator.trainable = False
10. #Building the generator
11. generator = build_generator()
12. #Defining the input for the generator
13. #and generating the images
14. z = Input(shape=(latent_dimensions,))
15. image = generator(z)
16. #Checking the validity of the generated image
17. valid = discriminator(image)
18. #Defining the combined model of the Generator and the Discriminator
19. combined_network = Model(z, valid)
20. combined_network.compile(loss='binary_crossentropy',
21.                           optimizer=Adam(0.0002,0.5))

```

Step 8: Training the network

```

1. num_epochs=15000
2. batch_size=32
3. display_interval=2500
4. losses=[]
5. #Normalizing the input
6. X = (X / 127.5) - 1.
7. #Defining the Adversarial ground truths
8. valid = np.ones((batch_size, 1))
9. #Adding some noise
10. valid += 0.05 * np.random.random(valid.shape)
11. fake = np.zeros((batch_size, 1))
12. fake += 0.05 * np.random.random(fake.shape)

```

```

13. for epoch in range(num_epochs):
14.     #Training the Discriminator
15.     #Sampling a random half of images
16.     index = np.random.randint(0, X.shape[0], batch_size)
17.     images = X[index]
18.     #Sampling noise and generating a batch of new images
19.     noise = np.random.normal(0, 1, (batch_size, latent_dimensions))
20.     generated_images = generator.predict(noise)
21.     #Training the discriminator to detect more accurately
22.     #whether a generated image is real or fake
23.     discm_loss_real = discriminator.train_on_batch(images, valid)
24.     discm_loss_fake = discriminator.train_on_batch(generated_images, fake)
25.     discm_loss = 0.5 * np.add(discm_loss_real, discm_loss_fake)
26.     #Training the Generator
27.     #Training the generator to generate images
28.     #which pass the authenticity test
29.     genr_loss = combined_network.train_on_batch(noise, valid)
30.     #Tracking the progress
31.     if epoch % display_interval == 0:
32.         display_images()

```

Output:

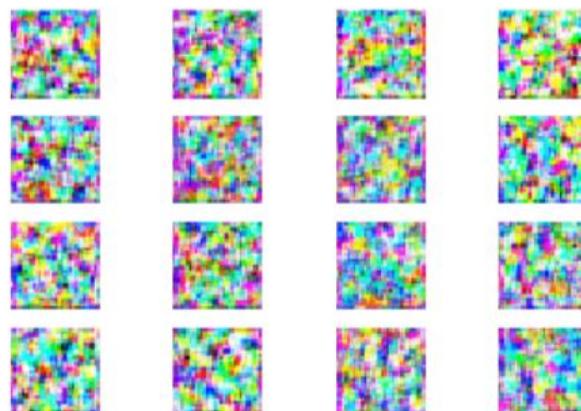


Fig1a.image generated at epoch size 20.

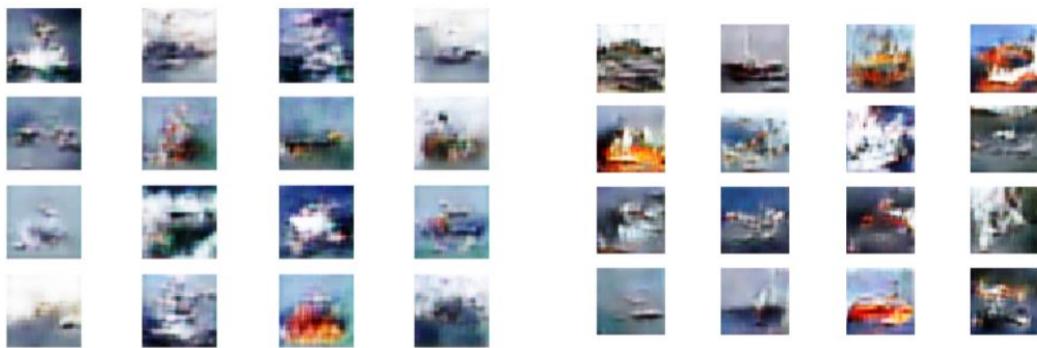


Fig1b image generated at epoch size 2000,

Fig1c.image generated at epoch size 5000